

Apache Kafka + Apache Spark

Apache Kafka + Apache Spark

- Apache Kafka and Apache Spark are part of the Data Science software stack.
- Sentiment Analysis uses Apache Kafka to store streaming data from Twitter and Apache Spark to process data. How does this work?
 - Collect tweets from Twitter and feed it into Apache Kafka
 - Kafka supplies this data in batches for processing in Apache Spark
 - Output the number of positive and negative words being mentioned



Sentiment Analysis using Kafka + Spark

- Step 1: Collect tweets from Twitter and feed it into Apache Kafka

```
import json
from kafka import SimpleProducer, KafkaClient
import tweepy
import configparser

# Note: Some of the imports are external python libraries. They are installed on the current machine.
# If you are running multinode cluster, you have to make sure that these libraries
# and current version of Python is installed on all the worker nodes.

class TweeterStreamListener(tweepy.StreamListener):
    """ A class to read the twitter stream and push it to Kafka """

    def __init__(self, api):
        self.api = api
        super(tweepy.StreamListener, self).__init__()
        client = KafkaClient("localhost:9092")
        self.producer = SimpleProducer(client, async = True,
                                      batch_send_every_n = 1000,
                                      batch_send_every_t = 10)

    def on_status(self, status):
        """ This method is called whenever new data arrives from live stream.
        We asynchronously push this data to kafka queue """
        msg = status.text.encode('utf-8')
        #print(msg)
        try:
            self.producer.send_messages(b'twitterstream', msg)
        except Exception as e:
            print(e)
            return False
        return True

    def on_error(self, status_code):
        print("Error received in kafka producer")
        return True # Don't kill the stream

    def on_timeout(self):
        return True # Don't kill the stream
```

Listens to the twitter stream continuously

Create a connection to Kafka that sends information in batches of 1000 tweets or every 10 seconds

Create a topic in Kafka under which tweets are saved

Either on encountering an error or timeout, do not stop the stream

Sentiment Analysis using Kafka + Spark

- Step 1: Collect tweets from Twitter and feed it into Apache Kafka

```
if __name__ == '__main__':  
  
    # Read the credentials from 'twitter.txt' file  
    config = configparser.ConfigParser()  
    config.read('twitter.txt')  
    consumer_key = config['DEFAULT']['consumerKey']  
    consumer_secret = config['DEFAULT']['consumerSecret']  
    access_key = config['DEFAULT']['accessToken']  
    access_secret = config['DEFAULT']['accessTokenSecret']  
  
    # Create Auth object  
    auth = tweepy.OAuthHandler(consumer_key, consumer_secret)  
    auth.set_access_token(access_key, access_secret)  
    api = tweepy.API(auth)  
  
    # Create stream and bind the listener to it  
    stream = tweepy.Stream(auth, listener = TweeterStreamListener(api))  
  
    # Custom Filter rules pull all traffic for those filters in real time.  
    # stream.filter(track = ['love', 'hate'], languages = ['en'])  
    stream.filter(locations=[-180,-90,180,90], languages = ['en'])
```

Create a Twitter developer account to get API access keys.

Authentication to connect to Twitter via Tweepy

Set the listener to the stream and filter information needed

Sentiment Analysis using Kafka + Spark

- Step 2: Kafka supplies this data in batches for processing in Apache Spark

```
def main():
    conf = SparkConf().setMaster("local[2]").setAppName("Streamer")
    sc = SparkContext(conf=conf)
    ssc = StreamingContext(sc, 10) # Create a streaming context with batch interval of 10 sec
    ssc.checkpoint("checkpoint")
    sc.setLogLevel("ERROR")

    pwords = load_wordlist("positive.txt")
    nwords = load_wordlist("negative.txt")

    counts = stream(ssc, pwords, nwords, 100)
    make_plot(counts)

def make_plot(counts):
    """
    Plot the counts for the positive and negative words for each timestep.
    Use plt.show() so that the plot will popup.
    """
    positive_count = []
    negative_count = []
```

Create Spark and Streaming context to stream in batches of 10 seconds

Call stream function to compute number of positive and negative words for 100 seconds

Your Code Goes Here

Sentiment Analysis using Kafka + Spark

- Step 2: Kafka supplies this data in batches for processing in Apache Spark

```
def stream(ssc, pwords, nwords, duration):
    kstream = KafkaUtils.createDirectStream(ssc, topics = ['twitterstream'], kafkaParams = {"metadata.broker.list": 'localhost:9092'})
    tweets = kstream.map(lambda x: x[1].encode("ascii", "ignore"))

    # Each element of tweets will be the text of a tweet.
    # You need to find the count of all the positive and negative words in these tweets.
    # Keep track of a running total counts and print this at every time step (use the pprint function).

    # Obtain list of words from tweets
    words = tweets.flatMap(lambda tweet: tweet.split())

    # Filter for words either in the positive list or negative list of words
    words = words.filter(lambda word: word in pwords or word in nwords)

    # Label and map each word to either 'positive' or 'negative', and 1 respectively
    words = words.map(lambda word: (word, 1))

    # Count and print the number of positive and negative words
    counts = words.reduceByKey(lambda a, b: a + b)

    # Let the counts variable hold the word counts for all time steps
    # You will need to use the foreachRDD function.
    # For our implementation, counts looked like:
    # [(["positive", 100), ("negative", 50)], [(["positive", 80), ("negative", 60)], ...]

    counts.foreachRDD(lambda rdd, _: pprint.pprint(rdd.collect()))

    # Start the computation
    ssc.start()
    ssc.awaitTerminationOrTimeout(duration)
    ssc.stop(stopGracefully=True)

    return counts
```

Connect to Kafka topic

Filter out words from tweet

Check if positive or negative

Map to 'positive' with count 1 or 'negative' with count 1

Roll up and get count

Continue for 100 seconds

Sentiment Analysis using Kafka + Spark

- Step 3: Output the number of positive and negative words being mentioned

```
-----  
Time: 2017-01-23 22:34:30  
-----
```

```
('positive', 145)  
('negative', 84)  
-----
```

```
Time: 2017-01-23 22:34:40  
-----
```

```
('positive', 285)  
('negative', 172)  
-----
```

```
Time: 2017-01-23 22:34:50  
-----
```

```
('positive', 460)  
('negative', 259)  
-----
```

```
Time: 2017-01-23 22:35:00  
-----
```

```
('positive', 616)  
('negative', 366)  
-----
```

```
Time: 2017-01-23 22:35:10  
-----
```

```
('positive', 759)  
('negative', 471)  
-----
```

Apache Spark + MongoDB

Apache Spark + MongoDB

- Apache Spark and MongoDB are part of the Data Science software stack.
- Word count uses Apache Spark to process data and the result is stored in MongoDB. How does this work?
 - Import data in text file
 - Apache Spark computes the word count
 - Result is stored in MongoDB collection



Word Count using Spark + MongoDB

- Step 1: Import data in text file. Apache Spark computes the word count.

```
# Import libraries
from pyspark import SparkContext, SparkConf
import json
from pyspark.sql.types import *
from pyspark.sql import SparkSession

# Create SparkSession
spark = SparkSession \
    .builder \
    .appName("Streamer") \
    .getOrCreate()

# Create SparkContext
sc = spark.sparkContext

# Import textfile and do word count
text_file = sc.textFile("./install-check/spark/textFile")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)

# Save result to output
counts.saveAsTextFile("./install-check/spark/output")

# Write to MongoDB
data = spark.read.json(counts)
data.write.format("com.mongodb.spark.sql.DefaultSource").mode("append").save()
data.show()
```

```
+-----+
|_corrupt_record|
+-----+
|('result', 1)|
|('textfile.', 1)|
|('to', 2)|
|('input', 1)|
|('the', 6)|
|('of', 1)|
|('has', 1)|
|('this', 1)|
|('located', 1)|
|('each', 1)|
|('currently', 1)|
|('coding.', 1)|
|('you', 1)|
|('count', 1)|
|('on', 1)|
|('check', 1)|
|('been', 1)|
|('Apache', 1)|
|('The', 3)|
|('by', 1)|
+-----+
only showing top 20 rows
```

Word Count using Spark + MongoDB

- Step 2: Result is stored in MongoDB collection.

```
1. sudo service mongod start
```

```
2. $SPARK_HOME/bin/spark-submit --conf "spark.mongodb.input.uri=mongodb://127.0.0.1/local.coll?readPreference=primaryPreferred" --conf "spark.mongodb.output.uri=mongodb://127.0.0.1/local.coll" --packages org.mongodb.spark:mongo-spark-connector_2.11:2.0.0 ./install-check/spark/testing.py
```

```
3. mongo
```

This will open the mongo shell. Type the following commands in the shell.

```
4. use local
```

```
5. db.coll.find().pretty()
```

You will see the results in MongoDB.

```
mycollection
> db.nongodbtest.find().pretty()
{
  "_id" : ObjectId("5a9cb0795ed30c1c631a72a7"),
  "_corrupt_record" : "('result', 1)"
}
{
  "_id" : ObjectId("5a9cb0795ed30c1c631a72a8"),
  "_corrupt_record" : "('textfile.', 1)"
}
{
  "_id" : ObjectId("5a9cb0795ed30c1c631a72a9"),
  "_corrupt_record" : "('to', 2)"
}
{
  "_id" : ObjectId("5a9cb0795ed30c1c631a72aa"),
```