
Deep Learning Parameter Optimization

Nagiza F. Samatova, samatova@csc.ncsu.edu
Professor, Department of Computer Science
North Carolina State University


Outline

- Model Hyperparameters
- Parameter Optimization Problem
- Parameter Optimization Strategies
 - Trial and error
 - Grid Search
 - Random Search
 - Bayesian Parameter Optimization
- Grid Search
- Random Search

Model Parameters and Hyperparameters

- **Model Parameters:**
 - Parameters that are learned during model training
 - Examples: weights, intercept/bias parameters
- **Hyperparameters:**
 - Parameters that **must be set** at the offset and are **not optimized during model training**
- **Examples of Hyperparameters:**
 - **The choice of a model:** deep feedforward neural network, random forest, support vector machines, etc.
 - **Model architecture:** the number of hidden layers, the number of nodes in each hidden layer, types of activation functions
 - **Model Convergence and Stopping Criteria:** the number of training epochs (iterations), threshold on the difference in optimization function between two subsequent iterations, learning rate
 - **Form of regularization to avoid model overfitting:** L_1 - or L_2 -norm regularizer, the number of dropout neurons,

Deep NN Models: Hyperparameters


```
38 system.time(ex1 <- h2o.deeplearning(  
39   x = xnames,  
40   y = "label",  
41   training_frame= h2odigits.train,  
42   validation_frame = h2odigits.test,  
43   activation = "RectifierWithDropout",  
44   hidden = c(100),  
45   epochs = 10,  
46   adaptive_rate = FALSE,  
47   rate = .001, ,  
48   input_dropout_ratio = 0,  
49   hidden_dropout_ratios = c(.2)  
50 ))
```

Validation Set Metrics:

=====

Extract validation frame with `h2o.getFrame`
MSE: (Extract with `h2o.mse`) 0.032
RMSE: (Extract with `h2o.rmse`) 0.18
Logloss: (Extract with `h2o.logloss`) 0.21
Mean Per-Class Error: 0.038
Confusion Matrix: Extract with `h2o.confusion`

Although ex1 took longer to train, it performs substantially better on the test data than does ex2.

```
52 system.time(ex2 <- h2o.deeplearning(  
53   x = xnames,  
54   y = "label",  
55   training_frame= h2odigits.train,  
56   validation_frame = h2odigits.test,  
57   activation = "RectifierWithDropout",  
58   hidden = c(100),  
59   epochs = 10,  
60   adaptive_rate = FALSE,  
61   rate = .01, ,  
62   input_dropout_ratio = 0,  
63   hidden_dropout_ratios = c(.2)  
64 ))
```

Validation Set Metrics:

=====

Extract validation frame with `h2o.getFrame`
MSE: (Extract with `h2o.mse`) 0.08
RMSE: (Extract with `h2o.rmse`) 0.28
Logloss: (Extract with `h2o.logloss`) 2.5
Mean Per-Class Error: 0.082
Confusion Matrix: Extract with `h2o.confusion`

Overarching Goal: Parameter Optimization

- **Motivation:** Some of the key reasons for poor model performance:
 - lack of the variables required for good prediction
 - not enough data to support training a complex enough model
 - **poorly tuned and optimized hyperparameters**
- **Goal:** To choose the best combination of and best values for the hyperparameters of the machine learning model.
 - **Assumption:** Better hyperparameters can often improve the accuracy of a model.

The values chosen for the hyperparameters can have a dramatic impact on the accuracy and training speed of a model.

Parameter Optimization Strategies

- Trial and error
- Grid Search:
 - pros: great if there are only a **few values** for a **few parameters**
 - cons: combinatorial (brute-force exhaustive enumeration of all possible combinations)
- Random Search:
 - pros: searching via random sampling; no need to pre-specify all the values to try and create all possible combinations
 - cons: computationally demanding for large sample sizes
- Bayesian Optimization (e.g., using **Spearmint Python library**):
 - iteratively adjusts a number of parameters so as to minimize some objective in as few runs as possible
 - <https://github.com/HIPS/Spearmint>

Grid Search for Optimal Parameters

- Basic Idea:
 - several values for hyperparameters are specified and
 - all possible combinations of these values are tried
- To create all possible combinations in R:
 - Use `expandGrid()` function in **gridExtra** package

DeepNeuralNetwork.ParameterOptimization.R

```
14 expand.grid(  
15   layers = c(1, 2, 4),  
16   epochs = c(50, 100),  
17   l1 = c(.001, .01, .05))
```

DeepNeuralNetwork.R

```
40 digits.m1 <- train(digits.X, digits.y,  
41   method = "nnet",  
42   tuneGrid = expand.grid(  
43     .size = c(5),  
44     .decay = 0.1),  
45   trControl = trainControl(method = "n  
46   MaxNWts = 10000,  
47   maxit = 100)
```

all possible combinations

| | layers | epochs | l1 |
|----|--------|--------|-------|
| 1 | 1 | 50 | 0.001 |
| 2 | 2 | 50 | 0.001 |
| 3 | 4 | 50 | 0.001 |
| 4 | 1 | 100 | 0.001 |
| 5 | 2 | 100 | 0.001 |
| 6 | 4 | 100 | 0.001 |
| 7 | 1 | 50 | 0.010 |
| 8 | 2 | 50 | 0.010 |
| 9 | 4 | 50 | 0.010 |
| 10 | 1 | 100 | 0.010 |
| 11 | 2 | 100 | 0.010 |
| 12 | 4 | 100 | 0.010 |
| 13 | 1 | 50 | 0.050 |
| 14 | 2 | 50 | 0.050 |
| 15 | 4 | 50 | 0.050 |
| 16 | 1 | 100 | 0.050 |
| 17 | 2 | 100 | 0.050 |
| 18 | 4 | 100 | 0.050 |

Random Search for Optimal Parameters

- What to specify for random sampling:
 - the values to randomly sample or
 - distributions to randomly draw from.
 - some limits, e.g., # of hidden layers in a range from 1 to 10
- Examples:
 - The **number of dropout neurons** will be drawn from **beta distribution** (20% dropouts in input layer and 50% on hidden layer): **dbeta()**
 - The depth or **number of layers** is sampled from 1 to 5
 - Adaptive learning rates (ρ and ϵ) are drawn from a **uniform distribution**: **runif()**
- How to do random sampling:
 - write a function that takes a seed and
 - then randomly samples a number of hyperparameters,
 - stores the sampled parameters,
 - runs the model, and
 - returns the results

Step-1: Set Up Random Search Parameters

Setting Up Model Hyperparameters

```
92 run <- function(seed,  
93                 name = paste0("m_", seed),  
94                 run = TRUE) {  
95   set.seed(seed)  
96  
97   p <- list(  
98     Name = name,  
99     seed = seed,  
100    depth = sample(1:5, 1),  
101    l1 = runif(1, 0, .01),  
102    l2 = runif(1, 0, .01),  
103    input_dropout = rbeta(1, 1, 12),  
104    rho = runif(1, .9, .999),  
105    epsilon = runif(1, 1e-10, 1e-4))  
106  
107   p$neurons <- sample(20:600, p$depth, TRUE)  
108   p$hidden_dropout <- rbeta(p$depth, 1.5, 1)/2
```

Step-2: Build Models for these Parameters

Building Models for Different Model Hyperparameters

```
110-   if (run) {  
111       model <- h2o.deeplearning(  
112         x = colnames(use.train.x),  
113         y = "Outcome",  
114         training_frame = h2oactivity.train,  
115         activation = "RectifierWithDropout",  
116         hidden = p$neurons,  
117         epochs = 100,  
118         loss = "CrossEntropy",  
119         input_dropout_ratio = p$input_dropout,  
120         hidden_dropout_ratios = p$hidden_dropout,  
121         l1 = p$l1,  
122         l2 = p$l2,  
123         rho = p$rho,  
124         epsilon = p$epsilon,  
125         export_weights_and_biases = TRUE,  
126         model_id = p$Name  
127     )
```

Step-3: Measure Models' Performance

Measure Model Performance for Different Model Hyperparameters

```
131     ## Measure performance on training data
132     p$MSE <- h2o.mse(model)
133     p$R2 <- h2o.r2(model)
134     p$Logloss <- h2o.logloss(model)
135     p$CM <- h2o.confusionMatrix(model)
136
137     ## Measure performance on testing data
138     perf <- h2o.performance(model, h2oactivity.test)
139     p$T.MSE <- h2o.mse(perf)
140     p$T.R2 <- h2o.r2(perf)
141     p$T.Logloss <- h2o.logloss(perf)
142     p$T.CM <- h2o.confusionMatrix(perf)
```

Step-4: Return Parameters, Model & Results

Return Model Results for Different Model Hyperparameters

```
148     return(list(  
149         Params = p,  
150         Model = model))
```

Step-5: Generate Results for Different Seeds

To make results reproducible, use different seeds:

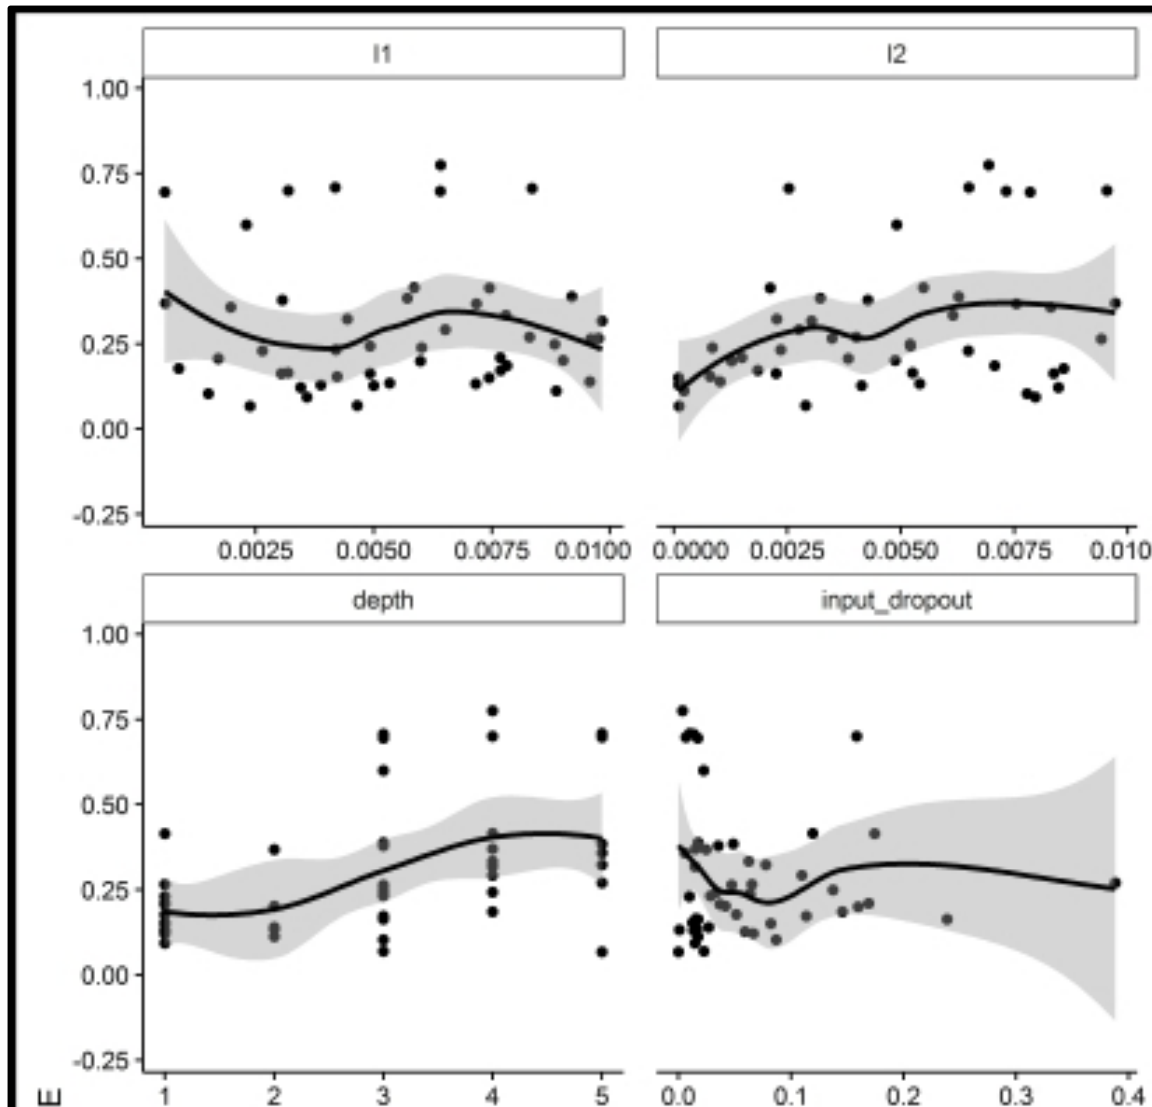
```
188 # To make the parameters reproducible,  
189 # specify a list of random seeds to loop through to run the models|  
190 use.seeds <- c(403L, 10L)  
191 # Step-5: To run the models simply by looping through the seeds:  
192 model.res <- lapply(use.seeds, run)
```

Step-6: Examine MSE for Diff. Parameters

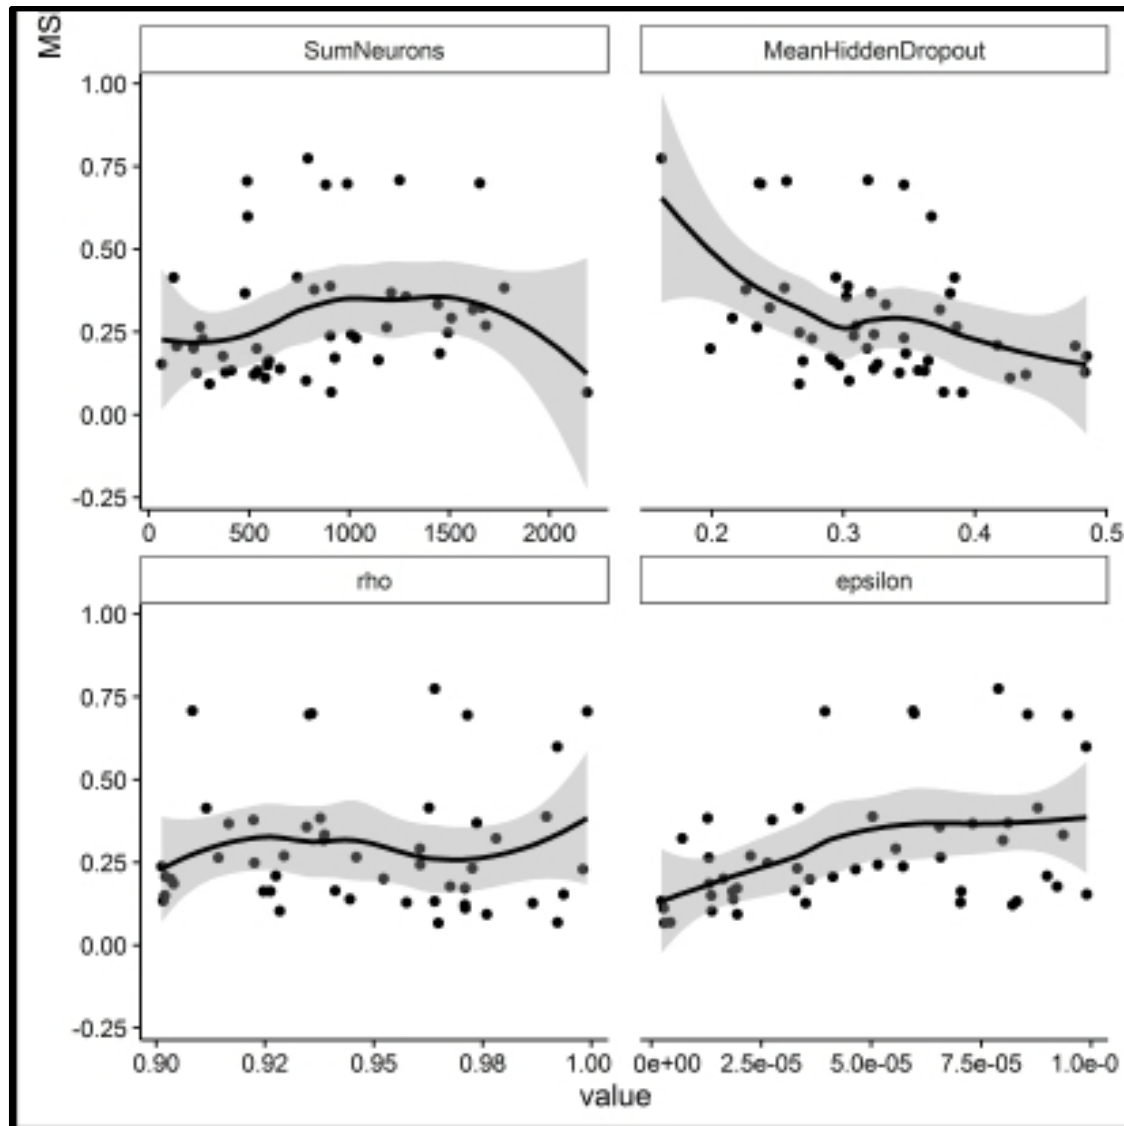
Plot the mean squared error (MSE) against the different parameters

```
196 model.res.dat <- do.call(rbind,  
197     lapply(model.res,  
198         function(x) with(x$params,  
199             data.frame(l1 = l1, l2 = l2,  
200                 depth = depth,  
201                 input_dropout = input_dropout,  
202                 sum_neurons = sum(neurons),  
203                 mean_hidden_dropout = mean(hidden_dropout),  
204                 rho = rho,  
205                 epsilon = epsilon,  
206                 MSE = T.MSE))))  
207  
208  
209 p.perf <- ggplot(melt(model.res.dat, id.vars = c("MSE")), aes(value, MSE)) +  
210     geom_point() +  
211     stat_smooth(colour = "black") +  
212     facet_wrap(~ variable, scales = "free_x", ncol = 2) +  
213     theme_classic()  
214  
215 print(p.perf)
```

Step-6: MSE Plots for Different Seeds



Step-6: MSE Plots for Different Seeds

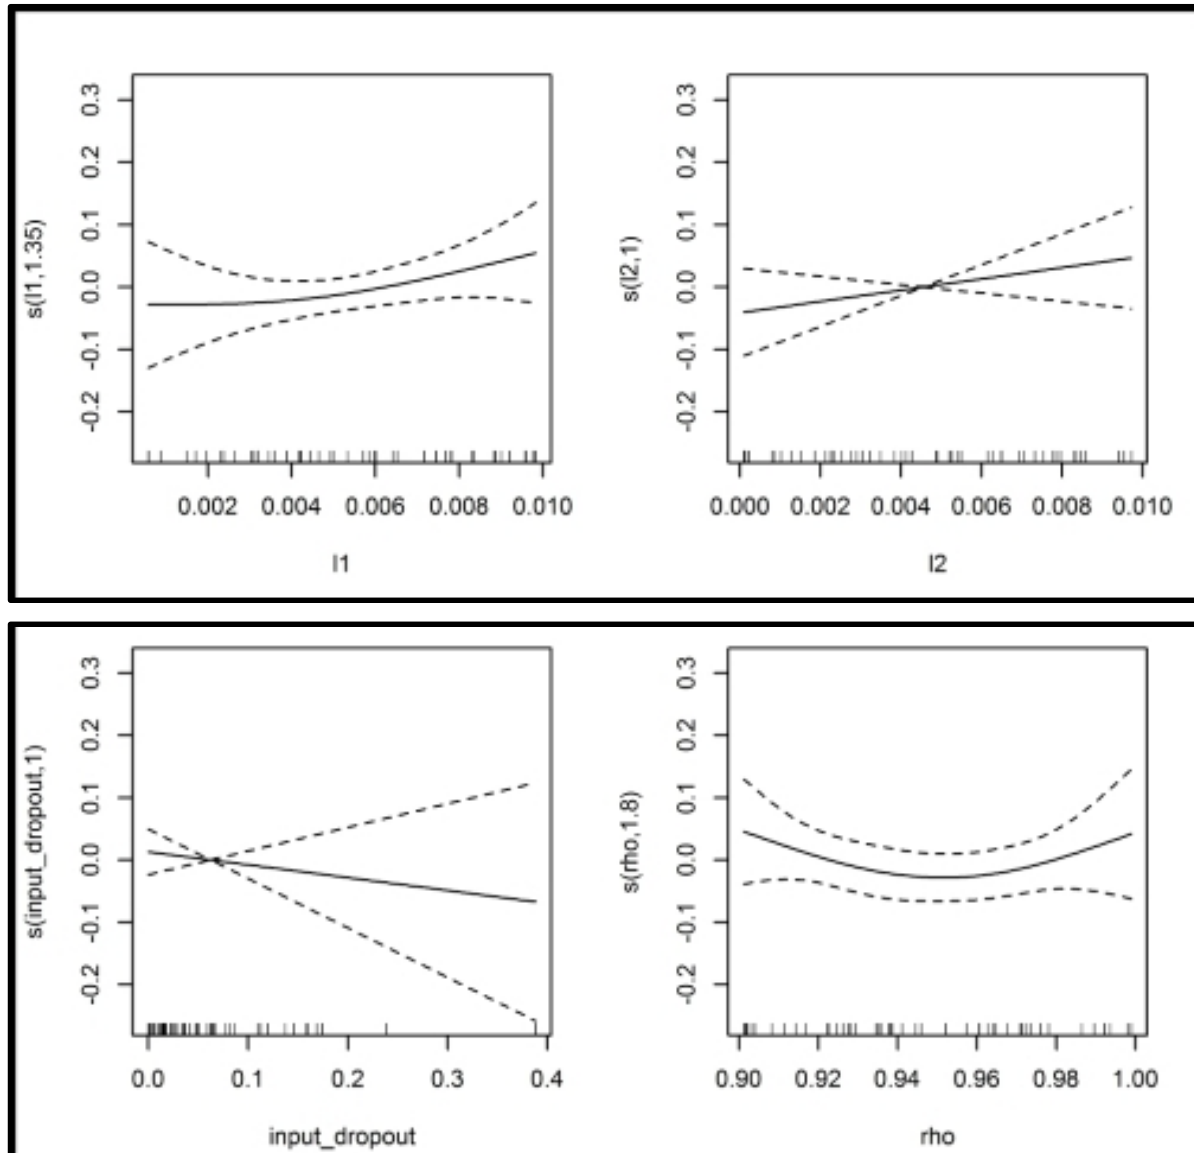


Step-7: Multivariate Parameter Relationships


Examine multivariate parameter relationships w.r.t. MSE

```
227 summary(m.gam <- gam(MSE ~ s(l1, k = 4) +
228                       s(l2, k = 4) +
229                       s(input_dropout) +
230                       s(rho, k = 4) +
231                       s(epsilon, k = 4) +
232                       s(MeanHiddenDropout, k = 4) +
233                       te(depth, SumNeurons, k = 4),
234                       data = model.res.dat))
235 for (i in 1:6) {
236   plot(m.gam, select = i)
237 }
238 plot(m.gam, select = 7)
```

Step-7: MSE Multivariate Relation Plots



Step-7: MSE Multivariate Relation Plots



Step-8: Choose Optimal Model Parameters

```
247 model.optimized <- h2o.deeplearning(  
248   x = colnames(use.train.x),  
249   y = "Outcome",  
250   training_frame = h2oactivity.train,  
251   activation = "RectifierWithDropout",  
252   hidden = c(500, 500, 500),  
253   epochs = 100,  
254   loss = "CrossEntropy",  
255   input_dropout_ratio = .08,  
256   hidden_dropout_ratios = c(.50, .50, .50),  
257   l1 = .002,  
258   l2 = 0,  
259   rho = .95,  
260   epsilon = 1e-10,  
261   export_weights_and_biases = TRUE,  
262   model_id = "optimized_model"  
263 )  
264  
265 h2o.performance(model.optimized, h2oactivity.test)  
266  
267 model.res.dat[which.min(model.res.dat$MSE), ]
```

References

- Chapter 19, Design and Analysis of Machine Learning Experiments
- Bengio, Y. (2012), **Section 3, Hyper-Parameters** in Practical Recommendations for Gradient-Based Training of Deep Architectures. In Neural Networks: Tricks of the Trade (pp. 437-478). Springer Berlin Heidelberg. (Also on the arXiv: <http://arxiv.org/pdf/1206.5533.pdf>)
- Zeiler, M. D. (2012), **ADADELTA**: An Adaptive Learning Rate Method. arXiv:1212.5701
- Bayesian optimization of hyperparameters, available online:
 - <https://github.com/HIPS/Spearmint>
 - see references on github for Spearmint library
- Joshua F. Wiley , R Deep Learning Essentials