

Word2Vec and Doc2Vec

Nagiza F. Samatova, samatova@csc.ncsu.edu

Professor, Department of Computer Science
North Carolina State University

Word2Vec is a tool to create vector space representations of words

- **“embeddings” is a term often used instead of “vector space representation”:**
 - Embedding means representation of an object embedded (placed) in a vector space.
- **Word2Vec produces vector space representations of words that:**
 - Capture co-occurrence of words in text
 - Captures semantic meaning by accounting for word order
- **Word2Vec is a shallow neural network**
 - Input layer, hidden layer, and output layer
 - The number of neurons in the hidden layer is the number of dimensions in the vector space
 - The neural net in Word2Vec is only used to obtain the weights, it is not used in the normal way – e.g. to perform classification with forward propagation.
- **Word2Vec has two different versions:**
 - continuous bag of words and
 - skip-gram

Ex: Semantic Meaning of Relationships in Vector Space

Test for linear relationships, examined by Mikolov et al. (2014)

a:b :: c:?



$$d = \arg \max_x \frac{(w_b - w_a + w_c)^T w_x}{\|w_b - w_a + w_c\|}$$

Word Analogies


man:woman :: king:?

+ king [0.30 0.70]

- man [0.20 0.20]

+ woman [0.60 0.30]

queen [0.70 0.80]



word2vec: Input, output, parameters, training...

● Input

- Corpus – the text that you train word2vec model on (e.g., the entire Wikipedia, your Inbox, all Harry Potter books, etc.)
- Vocabulary (it is extracted by word2vec but could be an input)

● Parameters

- Number of dimensions of the **target vector space** – N:
 - Usually small: 64, 128, ..., under 1,000
 - Unlike bag of words: ~100,000 of words, with each word as a dimension
- **context window** size:
 - Usually: 5-11 words
- others

● Output

- Vectors in the N-dimensional space (Euclidean space)
 - Each word is represented by an array of N numbers

Continuous Bag of Words (CBOW): One-word context

We assume that there is only one word considered per context.

Model will predict one target word given one context word

Neural network: One-word context

The input vector $x = \{x_1, \dots, x_V\}$ is a one-hot encoded vector

E.g. “eats” -> [0 0 0 1 0 0 0]


The weights between the input and the hidden layer can be represented by a $V \times N$ matrix W

Each row of W is the N -dimensional vector v_w

v_w is the vector representation of associated word of the input layer.

V is the vocabulary size

N is size of the hidden layer



The weights between the hidden and the output layer can be represented by a $V \times N$ matrix W'

CBOW: Bag of words context

We assume that there several words considered per context.

Model will predict one target word given many context word

context

target

Cookie monster eats → cookies



Continuous bag-of-words neural network architecture




Figure 2: Continuous bag-of-word model

Skip-Gram Model: One-word target -> context

We assume that there is only one word considered in target and many (window-size) in context.

Model will predict many context words given one target word

target

eats

context

Cookie monster ____ **cookies**



- Skip-gram example:
 - 3-grams:
 - Cookie, monster, cookies
 - Cookie,eats, cookies
 - monster, eats, cookies

Skip-Gram with Negative Sampling

Idea: don't just use 1's for the words you see in the context also use 0 for the words you **don't see** in the context

target

context

eats → Frobenius norm **eats** cookies

$$\|A\|_F = \left(\sum_{i=1}^m \sum_{j=1}^n a_{ij}^2 \right)^{1/2}$$

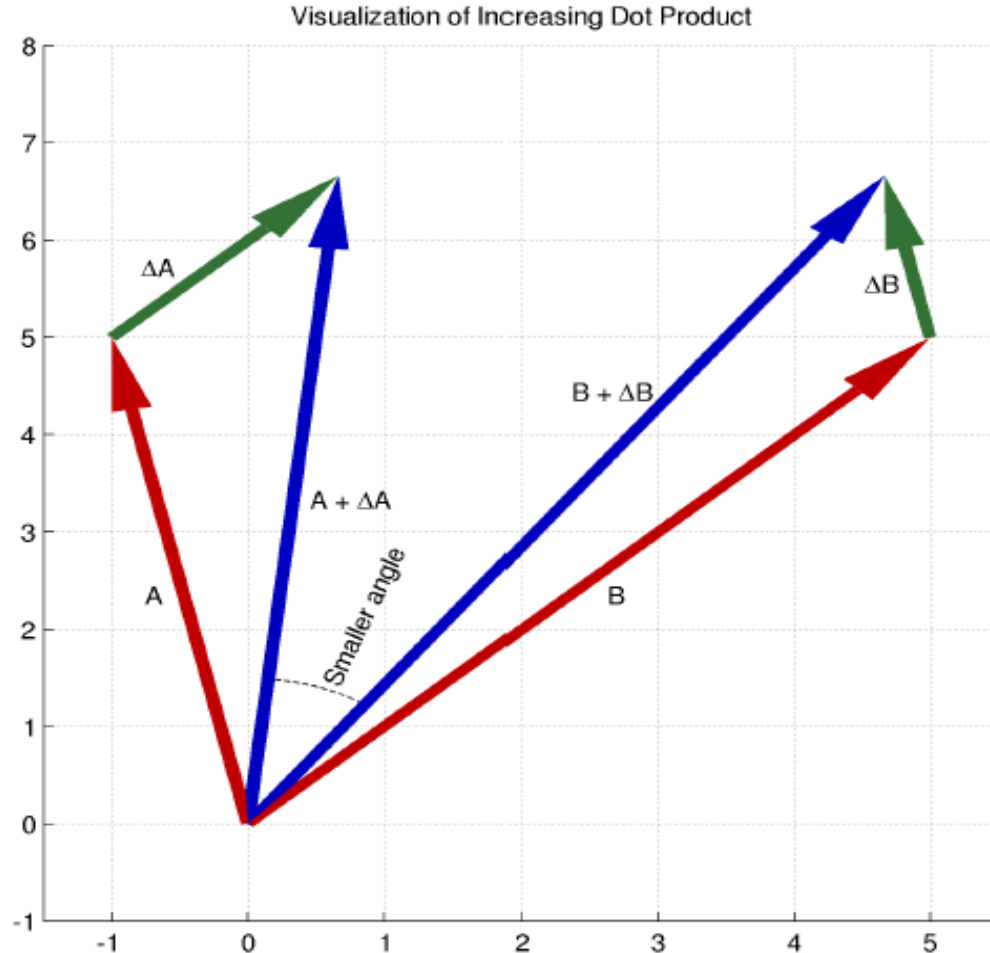


Skip-gram overview

- **Each word corresponds to two vectors, one is the target word and the other as a context word.**
- **The algorithm starts with random vectors for the context and zero vectors for the target words.**
- **Each time one word occurs in the context of the other, the context vector and target vector are modified slightly so that their product is slightly larger.**

Dissertation by Eric D. Moyer
“What machines understand
about personality words after
reading the news”

Intuition: minimization of context-target dot-product



- Each time one word occurs in the context of the other, the context vector and target vector are modified slightly so that their dot product is slightly larger.
- the dot product between a word and a target will be proportional to the probability of that word appearing in the context of that target when compared to the dot products of other words also appearing in that context.
- This means that words will cluster around their most common contexts.

Training skip-gram model on a tiny example

Corpus: Mary had a little lamb, little
lamb little, little lamb Mary had a
little lamb, little lamb little, little lamb
Mary had a little lamb, little lamb
little, little lamb Mary had a little
lamb, little lamb little, little lamb

Vocabulary ($V=5$): Mary, had, a,
little, lamb

Dimensionality: $N = 2$

Context window size: $w = 1$

Training skip-gram model

Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb

Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb

Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb
Mary had a little lamb, little lamb, little lamb, little lamb


context

target

One-hot encoding

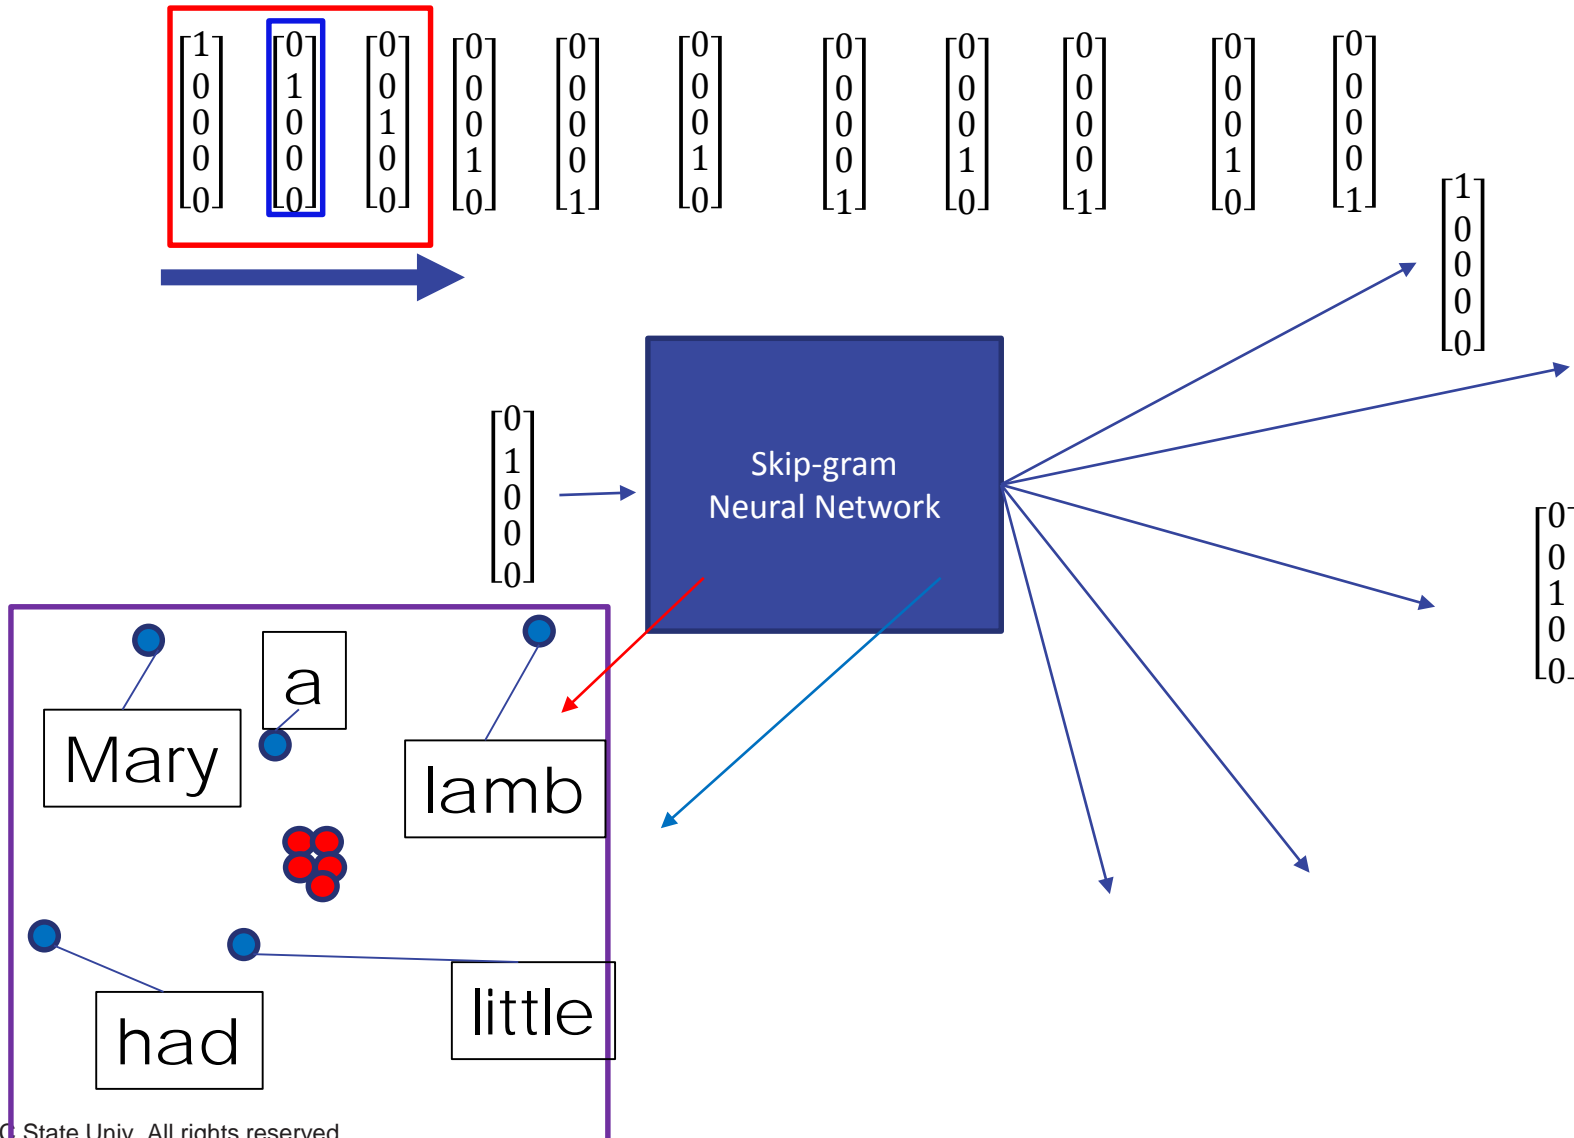
| Index | Word | One-hot Encoding vector |
|-------|--------|-------------------------|
| 0 | Mary | [1,0,0,0,0] |
| 1 | had | [0,1,0,0,0] |
| 2 | a | [0,0,1,0,0] |
| 3 | little | [0,0,0,1,0] |
| 4 | lamb | [0,0,0,0,1] |

Mary had a little lamb, little lamb, little lamb, little lamb



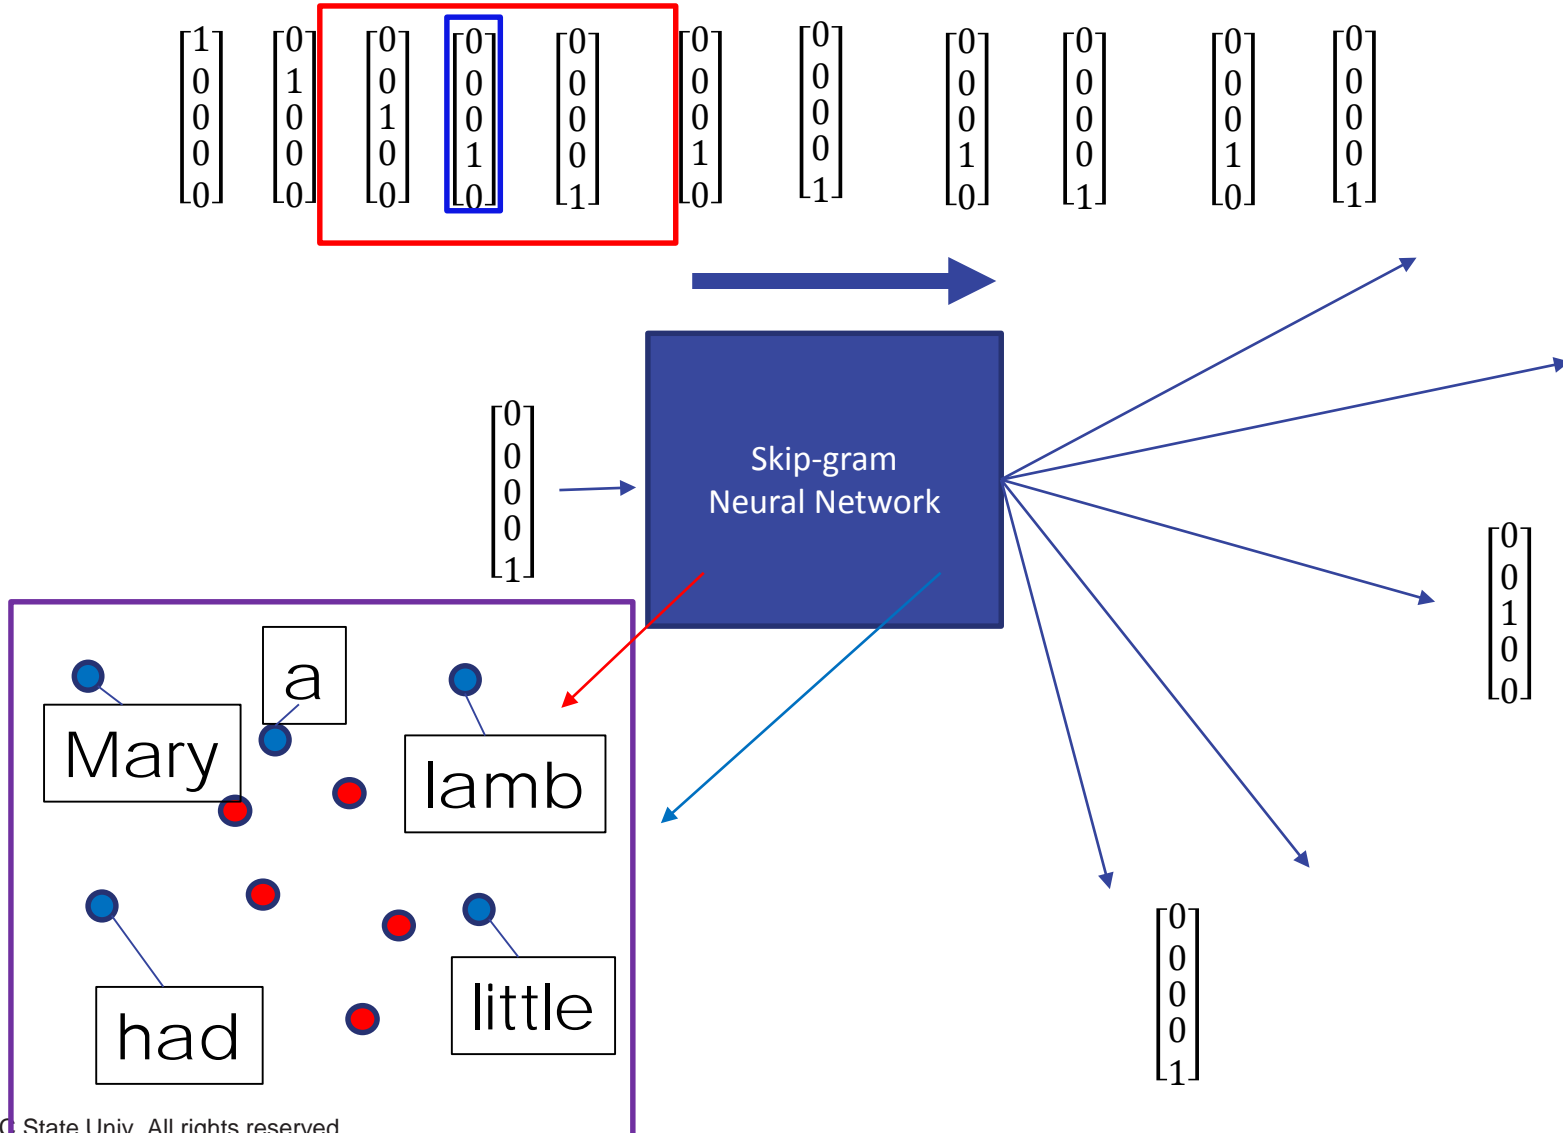
Initialized weights in neural net and start

Mary had a little lamb, little lamb, little lamb, little lamb



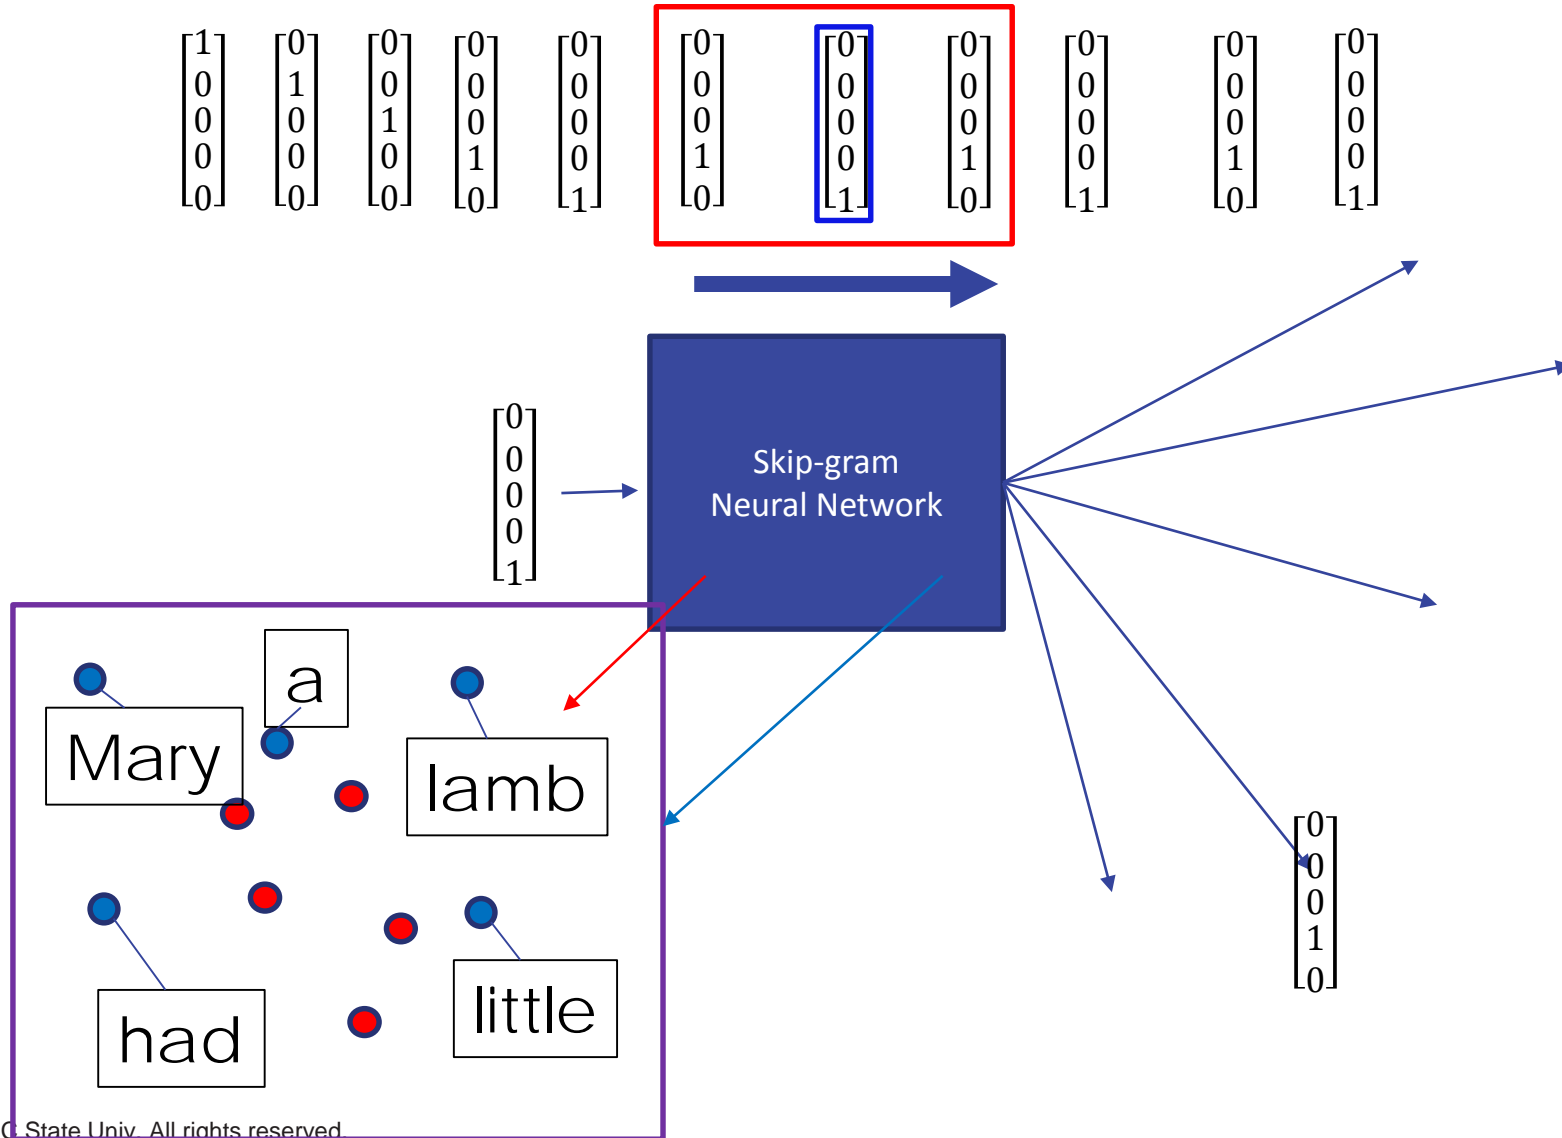
Training neural network

Mary had a little lamb, little lamb, little lamb, little lamb



Training neural network

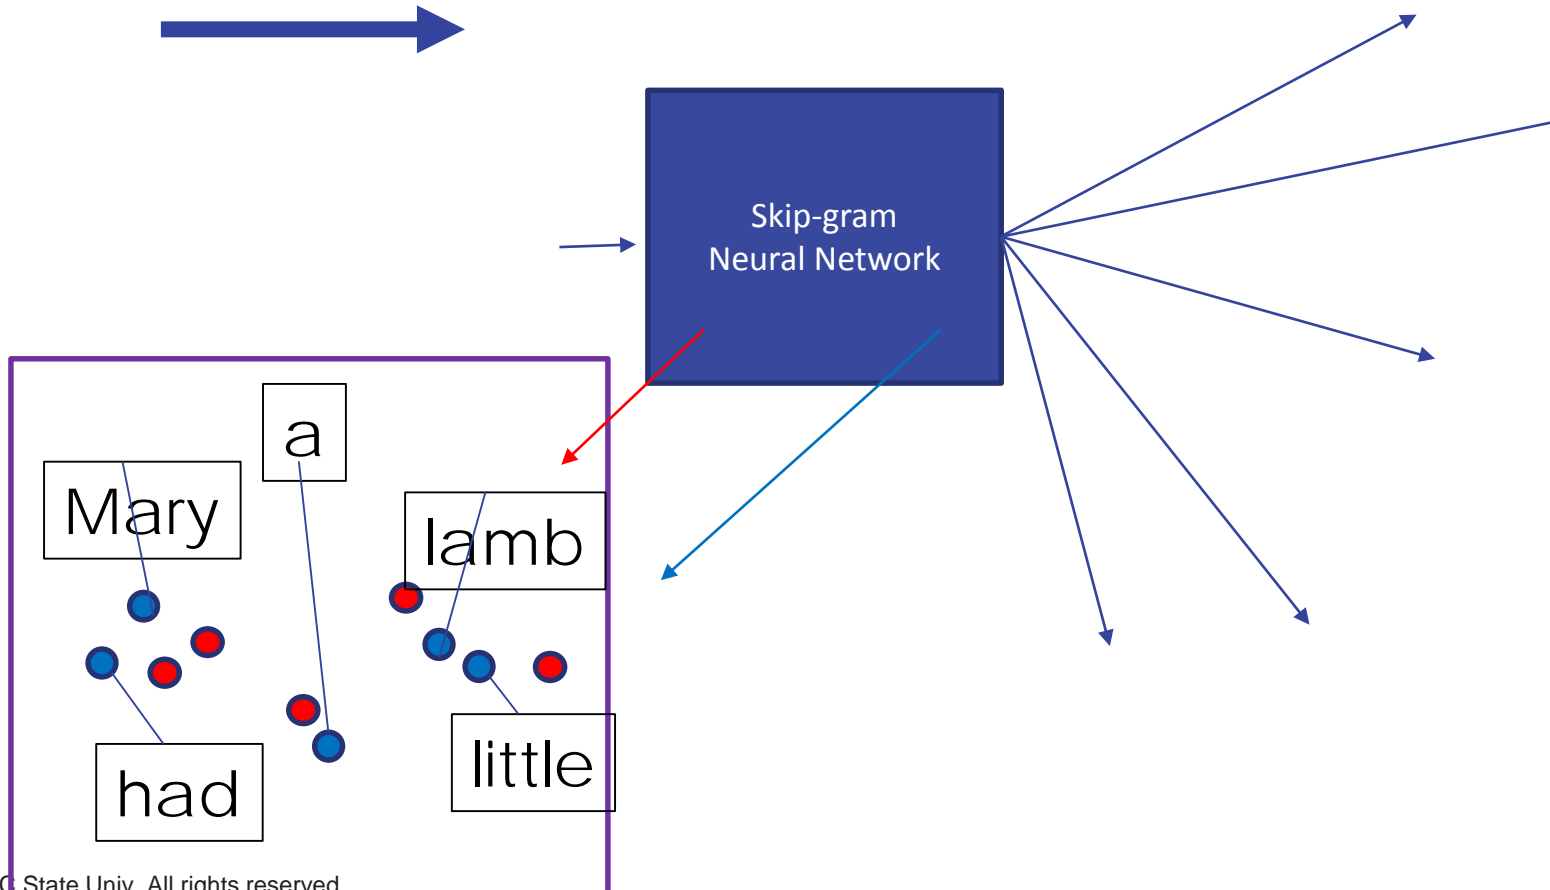

Mary had a little lamb, little lamb, little lamb, little lamb



Finished training neural network

Mary had a little lamb, little lamb, little lamb, little lamb

| | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|
| $\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$ |
|---|---|---|---|---|---|---|---|---|---|---|



Skip-gram model neural net

- The **target** word is in the **input** layer.
- The **context** words are in the **output** layer

v_{w_I} is the input vector of the only word on the input layer

For a given input word, and assuming its index in the vocabulary k , one-hot input vector representation is x .

$$x_i = \begin{cases} 1 & \text{for } i = k \\ 0 & \text{for } i \neq k \end{cases}$$
$$v_{w_I} := h = x^T W = W_{(k, \cdot)}$$

It is the k -th row of W , which is the input->hidden weight matrix




Figure 3: The skip-gram model.


Skip-gram model: details

On the output layer we output C multinomial distributions

Each output computed using the same hidden->output weight matrix W' and softmax link function:

$$y_{c,j} = \frac{\exp(u_{c,j})}{\sum_{j'=1}^V \exp(u_{j'})}$$

- Also note that we hope that:
- $y_{c,j} = p(w_{c,j} = w_{O,c} | w_I)$
- $w_{c,j}$ is the j -th word on the c -th panel of the output layer
- $w_{O,c}$ is the actual c -th word in the output context words
- conditional probability of observing context word given a target word




Skip-gram model: output

$$y_{c,j} = \frac{\exp(u_{c,j})}{\sum_{j'=1}^V \exp(u_{j'})}$$

$$u_j = \mathbf{v}_{w_j}'^T \cdot \mathbf{h} \text{ for } c = 1, 2, \dots, C$$

- $\mathbf{v}_{w_j}'^T$ is the output vector of the j -th word in the vocabulary, w_j
- $\mathbf{v}_{w_j}'^T$ is a column from the hidden->output matrix, \mathbf{W}'



Python implementation of word2vec

- <https://radimrehurek.com/gensim/models/word2vec.html>
<https://rare-technologies.com/word2vec-tutorial/>

```
>>> model = Word2Vec(sentences, size=100, window=5, min_count=5, workers=4)
```

```
>>> model.most_similar(positive=['woman', 'king'], negative=['man'])  
[('queen', 0.50882536), ...]
```

```
>>> model.doesnt_match("breakfast cereal dinner lunch".split())  
'cereal'
```

```
>>> model.similarity('woman', 'man')  
0.73723527
```

```
>>> model['computer'] # raw numpy vector of a word  
array([-0.00449447, -0.00310097,  0.02421786, ...], dtype=float32)
```

Doc2Vec: Distributed Representations of Sentences and Documents

- Doc2Vec has 2 models:
 - Paragraph Vector Distributed Bag of Words (PV-DBOW)
 - Paragraph Vector Distributed Memory (PV-DM)

https://cs.stanford.edu/~quocle/paragraph_vector.pdf

Distributed Representations of Sentences and Documents

Quoc Le
Tomas Mikolov

Google Inc, 1600 Amphitheatre Parkway, Mountain View, CA 94043

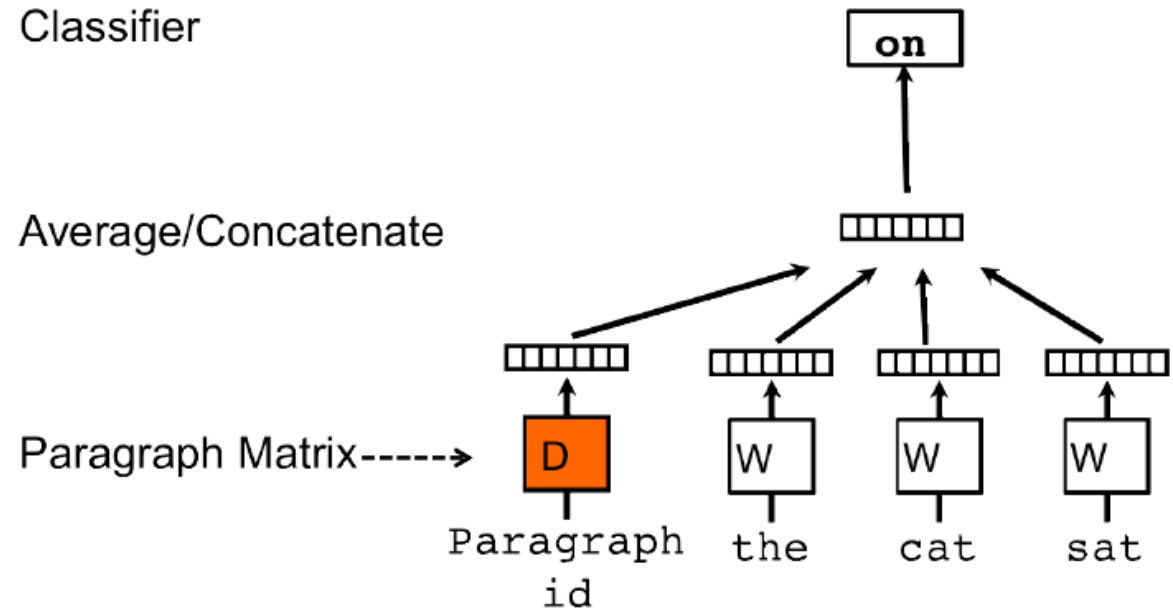
QVL@GOOGLE.COM
TMIKOLOV@GOOGLE.COM

Doc2Vec

- “In our model, the vector representation is trained to be useful for predicting words in a paragraph.
- ...we concatenate the paragraph vector with several word vectors from a paragraph and predict the following word in the given context.
- Both word vectors and paragraph vectors are trained by the stochastic gradient descent and backpropagation.
- While paragraph vectors are unique among paragraphs, the word vectors are shared.
- At prediction time, the paragraph vectors are inferred by fixing the word vectors and training the new paragraph vector until convergence.”

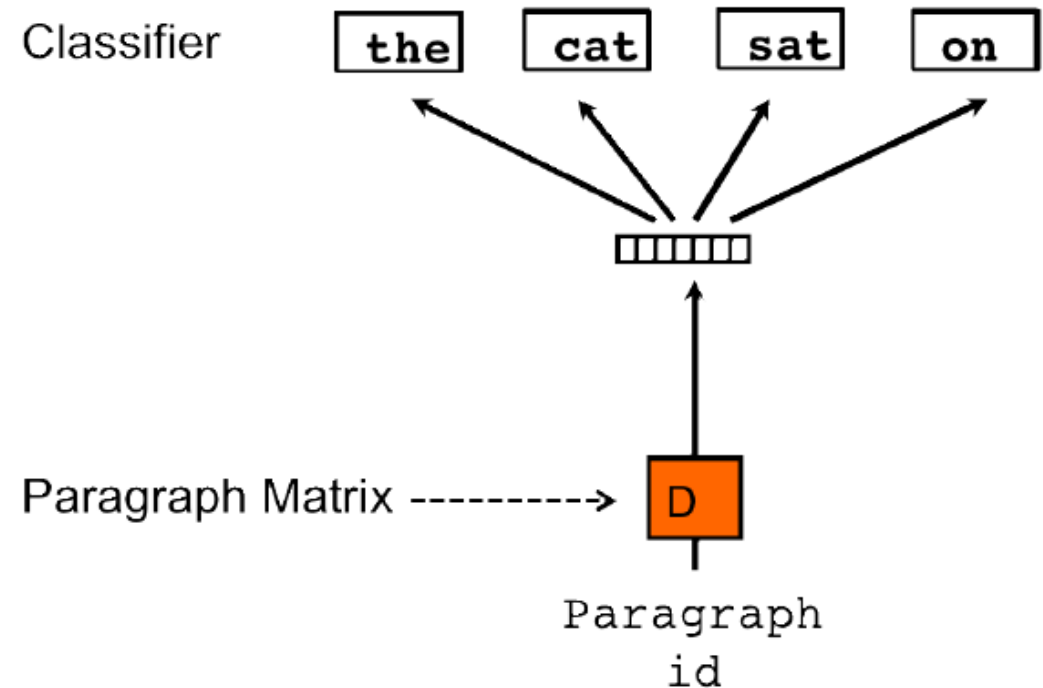
Paragraph vector: distributed memory

- every paragraph is mapped to a unique **vector**, represented by a column in matrix **D**
- and every word is also mapped to a unique **vector**, represented by a column in matrix **W**.
- The paragraph vector and word vectors are averaged or concatenated to predict the next word in a context.



Paragraph vector: distributed bag of words

- ...ignore the context words in the input, but force the model to predict words randomly sampled from the paragraph in the output.
- In reality, what this means is that at each iteration of stochastic gradient descent, we sample a text window, then sample a random word from the text window and form a classification task given the Paragraph Vector.



Doc2Vec usage with gensim package in Python

<https://radimrehurek.com/gensim/models/doc2v>

```
>>> model = Doc2Vec(documents, size=100, window=8, min_count=5, workers=4)
```

```
>>> trained_model.n_similarity(['sushi', 'shop'], ['japanese', 'restaurant'])  
0.61540466561049689
```

```
>>> trained_model.n_similarity(['restaurant', 'japanese'], ['japanese', 'restaurant'])  
1.0000000000000004
```

- `class gensim.models.doc2vec.Doc2Vec(documents=None, size=300, alpha=0.025, window=8, min_count=5, max_vocab_size=None, sample=0, seed=1, workers=1, min_alpha=0.0001, dm=1, hs=1, negative=0, dbow_words=0, dm_mean=0, dm_concat=0, dm_tag_count=1, docvecs=None, docvecs_mapfile=None, comment=None, trim_rule=None, **kwargs)`

Project: Sentiment analysis with Doc2Vec

- **perform sentiment analysis over IMDB movie reviews and Twitter data.**
- **build BOW, Doc2Vec, Word2Vec models with labeled training and testing data to evaluate the model.**
- **classify tweets or movie reviews as either positive or negative given labeled training** For classification, you will experiment with logistic regression as well as a Naive Bayes classifier from python's well-regarded machine learning package scikit-learn.

Project steps

1. Load datasets

```
34 def main():  
35     train_pos, train_neg, test_pos, test_neg = load_data(path_to_data)  
36
```

2. Train Doc2Vec model

```
62 if method == "d2v":  
63     train_pos_vec, train_neg_vec, test_pos_vec, test_neg_vec =  
64     feature_vecs_DOC(train_pos, train_neg, test_pos, test_neg)  
    #filename = './'+path_to_data+'train_pos_vec_d2v.txt'
```

3. Train logistic regression model with Doc2Vec vector-features

```
77     nb_model, lr_model = build_models_DOC(train_pos_vec,  
        train_neg_vec)
```

4. Evaluate the model on test data

```
113  
114     evaluate_model(lr_model, test_pos_vec, test_neg_vec, True)  
115
```

Training Doc2Vec steps:

Step 1. Create TaggedDocument-s from lists of words

```
197 labeled_train_pos = [TaggedDocument(words, ["TRAIN_POS_" + str(i)])  
    for i, words in enumerate(train_pos)]  
198 labeled_train_neg = [TaggedDocument(words, ["TRAIN_NEG_" + str(i)])  
    for i, words in enumerate(train_neg)]  
199 labeled_test_pos = [TaggedDocument(words, ["TEST_POS_" + str(i)])  
    for i, words in enumerate(test_pos)]  
200 labeled_test_neg = [TaggedDocument(words, ["TEST_NEG_" + str(i)])  
    for i, words in enumerate(test_neg)]
```

Step 2. Initialize the model

```
203 model = Doc2Vec(min_count=1, window=10, size=100, sample=1e-4,  
    negative=5, workers=4)
```

```
205 sentences = labeled_train_pos + labeled_train_neg + labeled_test_pos  
    + labeled_test_neg  
206 model.build_vocab(sentences)
```

Step 3. Train the model

```
210 for i in range(5):  
211     print("Training iteration %d" % (i))  
212     random.shuffle(sentences)  
213     model.train(sentences, total_examples=model.corpus_count, epochs=  
        model.iter)
```

Training Doc2Vec steps:

Step 4. Extract vectors

```
216 # Use the docvecs function to extract the feature vectors for the
217 training and test data
218 train_pos_vec = [model.docvecs["TRAIN_POS_" + str(i)] for i in range
219 (len(labeled_train_pos))]
220 train_neg_vec = [model.docvecs["TRAIN_NEG_" + str(i)] for i in range
221 (len(labeled_train_neg))]
222 test_pos_vec = [model.docvecs["TEST_POS_" + str(i)] for i in range(
223 len(labeled_test_pos))]
224 test_neg_vec = [model.docvecs["TEST_NEG_" + str(i)] for i in range(
225 len(labeled_test_neg))]
226
227 # Return the four feature vectors
228 return train_pos_vec, train_neg_vec, test_pos_vec, test_neg_vec
```


Build logistic regression model with Doc2Vec vectors

```
342 def build_models_DOC(train_pos_vec, train_neg_vec):
343     """
344     Returns a GaussianNB and LogisticRegression Model that are fit to the
345     training data.
346     """
347     Y = ["pos"]*len(train_pos_vec) + ["neg"]*len(train_neg_vec)
348     # Use sklearn's GaussianNB and LogisticRegression functions to fit
349     # two models to the training data.
350     # For LogisticRegression, pass no parameters
351     X = train_pos_vec + train_neg_vec
352     nb_model = sklearn.naive_bayes.GaussianNB()
353     nb_model.fit(X, Y)
354     lr_model = sklearn.linear_model.LogisticRegression()
355     lr_model.fit(X, Y)
356     return nb_model, lr_model
```

Evaluating the model

```
Doc2Vec
Training iteration 0
Training iteration 1
Training iteration 2
Training iteration 3
Training iteration 4
end of training
Naive Bayes
-----
predicted:      pos      neg
actual:
pos             5482     7018
neg             2262     10238
accuracy: 0.628800

Logistic Regression
-----
predicted:      pos      neg
actual:
pos             10745     1755
neg             1700     10800
accuracy: 0.861800
```

```
def evaluate_model(model, test_pos_vec, test_neg_vec, print_confusion=
False):
    """
    Prints the confusion matrix and accuracy of the model.
    """
    # Use the predict function and calculate the true/false positives
    and true/false negative.
    pos_results = list(model.predict(test_pos_vec))
    neg_results = list(model.predict(test_neg_vec))
    tp = pos_results.count("pos")
    tn = neg_results.count("neg")
    fn = pos_results.count("neg")
    fp = neg_results.count("pos")
    accuracy = float(tp + tn) / (tp + tn + fp + fn)

    if print_confusion:
        print("predicted:\tpos\tneg")
        print("actual:")
        print("pos\t\t%d\t%d" % (tp, fn))
        print("neg\t\t%d\t%d" % (fp, tn))
    print("accuracy: %f" % (accuracy))
```

**If you would like to write your own
implementation of word2vec's neural net ...**

Error function for Skip-gram model

$$\begin{aligned} E = -\log p(w_{O,1}, w_{O,2}, \dots, w_{O,C} | w_I) = \\ -\log \prod_{c=1}^C \frac{\exp(u_{c,j_c^*})}{\sum_{j'=1}^V \exp(u_{j'})} = \\ -\sum_{c=1}^C u_{j_c^*} + C \cdot \log \sum_{j'=1}^V \exp(u_{j'}) \end{aligned}$$

j_c^* is the index of the actual c -th output context word in the vocabulary

Now we need the gradient for the stochastic gradient descent

Prediction errors and gradients

Taking the derivative of E with regard to the net input of every node on every panel of the output layer, $u_{c,j}$ and setting it as prediction error:

Defining El_j as a sum of prediction errors over all context words:

Taking the derivative of E with regard to the hidden->output matrix W' with elements w'_{ij}

Thus we obtain the update equation for the hidden->output matrix W'

$$e_{c,j} = \frac{\partial E}{\partial u_{c,j}} = y_{c,j} - t_{c,j}$$

$$El_j = \sum_{c=1}^C e_{c,j}$$

$$\begin{aligned} \frac{\partial E}{\partial w'_{ij}} &= \sum_{c=1}^C \frac{\partial E}{\partial u_{c,j}} \cdot \frac{\partial u_{c,j}}{\partial w'_{ij}} = El_j \cdot h_i \\ w'_{ij}^{new} &= w'_{ij}^{old} - \eta \cdot El_j \cdot h_i \end{aligned}$$

Update Equations for Skip-gram Model

$$\mathbf{v}'_{w_j}{}^{new} = \mathbf{v}'_{w_j}{}^{old} - \eta \cdot EI_j \cdot \mathbf{h}$$

update of “context” vectors

$$\mathbf{v}_{w_I}^{new} = \mathbf{v}_{w_I}^{old} - \eta \cdot \mathbf{EH}$$

update of “target” vectors

$$EH_i = \sum_{j=1}^V EI_j \cdot w'_{ij}$$

$e_{c,j} := \hat{y}_{c,j}^1 - t_{c,j}$ is the **prediction error** on the node

$\mathbf{EI} = \{EI_1, \dots, EI_V\}$ is a V –dimensional vector, the **sum of all prediction errors** over all context words

$$EI_j = \sum_{c=1}^C e_{c,j}$$