

---

# Time Series (TS) Forecasting

## Differencing & Smoothing

**Nagiza F. Samatova**, [samatova@csc.ncsu.edu](mailto:samatova@csc.ncsu.edu)

Professor, Department of Computer Science  
North Carolina State University

Senior Scientist, Computer Science & Mathematics Division  
Oak Ridge National Laboratory

# Outline

---

- **Data-driven vs. Model-based Methods**
  - **Local vs. Global Patterns**
- **Differencing Methods**
  - **De-trending**
  - **De-seasonalizing**
- **Smoothing Methods**
  - **Moving Averages**
  - **Exponential Smoothing**
  - **ets() function for TS forecasting**

# Basic Notation


Symbol	Definition
$t = 1, 2, 3, \dots,$	An index for the time period of interest; e.g., for a <i>daily</i> time period, $t = 1$ means day 1, $t = 2$ means day 2, etc.
$y_1, y_2, \dots, y_T$	A series of $T$ values measure over $T$ time periods; e.g., for the annual average stock price, $y_1$ denotes the price for year 1, $y_2$ denotes the price for year 2, etc.
$F_t$ or $\hat{y}_t$	The forecast value for time period $t$
$F_{t+k}$ or $\widehat{y_{t+k}}$	The $k$ -step-ahead forecast when forecasting time is $t$ ; e.g., $F_{t+1}$ is the forecast for time period $(t + 1)$ made during the time period $t$
$e_t = y_t - F_t$	The forecast error for time period $t$

# TS Parts: **Systematic** vs **Non-systematic**

TS Part	Definition	Detection	How to deal w/
<b>Level</b>	Average value of ts		
<b>Trend</b>	Long-term increase decrease in the data	lag.plot	De-trend via lag-1 differencing
<b>Seasonality</b>	Variations occurring during known periods of the year (monthly, quarterly, holidays)	lag.plot, Acf plots	De-seasonalize via lag-k differencing
<b>Cycles</b>	Other oscillating patterns about the trend (e.g., business or economic conditions)		
<b>Auto-correlation</b>	Correlation between neighboring points in ts	Acf, lag.plot	
<b>Noise</b>	Residuals after level, trend, seasonality, and cycles are removed	Normality tests	

# Monthly Housing Sales


```
plot(hsales, ylab="Monthly housing sales (millions)")
```



- Strong seasonality within each year
- Cyclic behavior 6-10 years
- No trend over this period

# US Treasury Bill Contracts


```
plot(ustreas, ylab="US treasury bill contracts")
```



- No seasonality
- Downward trend
- No cyclic behavior over this period

# Monthly Electricity Production


```
plot(elec, ylab="Monthly Electricity Production")
```



- Strong seasonality
- Increasing trend
- No cyclic behavior over this period

# Daily Changes in Dow-Jones Index

```
plot(diff(dj,1), ylab="Daily Changes in Dow-Jones Index")
```



- No seasonality
- No trend
- No cyclic behavior over this period
- Random, unpredictable fluctuations, like white noise



# Additive and Multiplicative TS Components

---

- A time series with **additive** components can be modeled as:

$$y_t = \textit{Level} + \textit{Trend/Cycles} + \textit{Seasonality} + \textit{Noise}$$

- A time series with **multiplicative** components is modeled as:

$$y_t = \textit{Level} \times \textit{Trend/Cycles} \times \textit{Seasonality} \times \textit{Noise}$$

- Forecasting methods attempt to isolate the systematic part and quantify the noise level.
  - The systematic part is used for generating point forecasts
  - The level of noise helps assess the uncertainty associated with the point forecasts

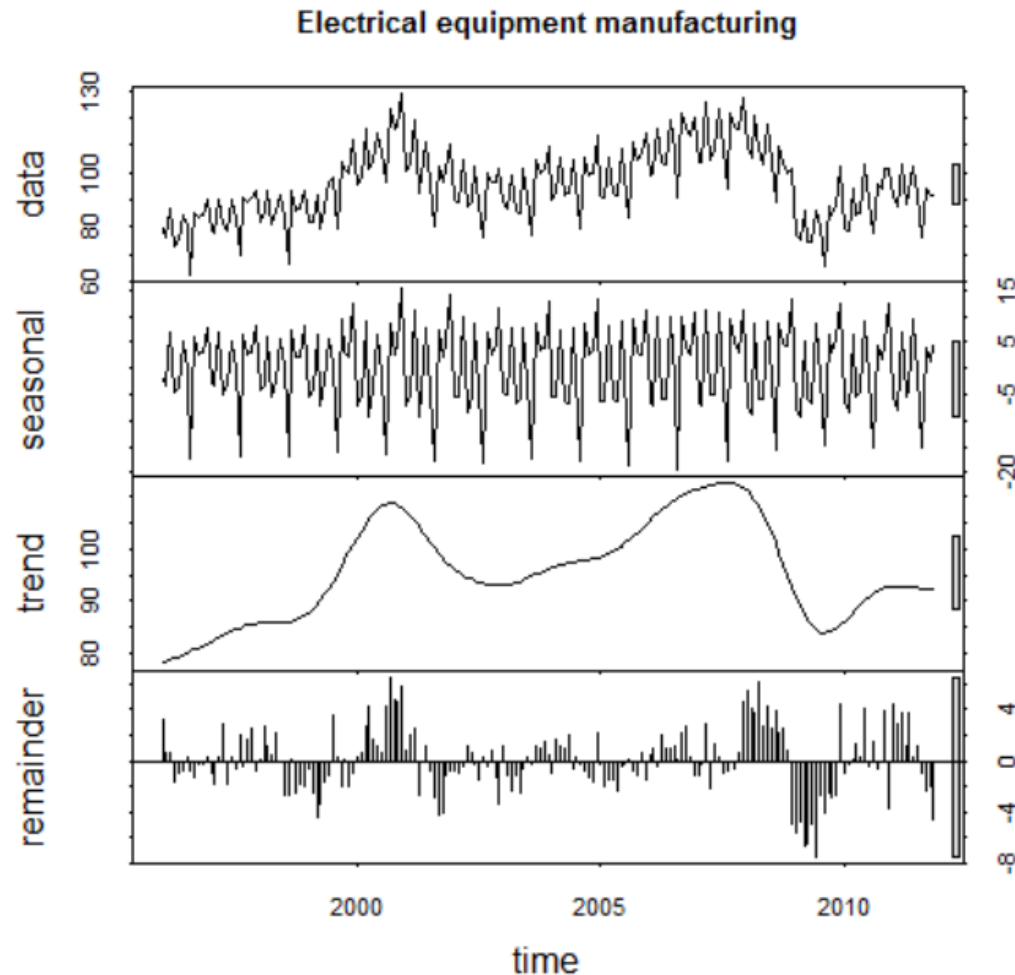
# Additive vs Multiplicative Decomposition

---

- **Additive model** is appropriate if the magnitude of the seasonal fluctuations or **the variation around the trend-cycle does NOT vary with the level of the time series**:
  - `decompose()`
  - `stl()`: for time series with additive seasonality
- **Multiplicative model** is appropriate if the magnitude of the seasonal fluctuations or **the variation around the trend-cycle appears to be proportional to the level of the time series**
  - Rather than building a multiplicative model, transform the time series to stabilize the variance over time and then use the additive model

# Example: TS Decomposition with STL

```
fit <- stl ( elecequip, s.window=5 )  
plot(fit, main="Electrical equipment manufacturing")
```




# Seasonality- vs. Trend-adjusted Time Series

---

- If the variation due to seasonality is not of primary interest then generate seasonally adjusted time series for subsequent analysis:
  - Additive model:  $y_t - S_t$ , where  $S_t$  is the additive seasonal component
  - Multiplicative model:  $y_t/S_t$ , where  $S_t$  is the multiplicative seasonal component
- If the purpose is to analyze and interpret turning points, then it is better to use trend-cycle component rather than seasonally adjusted data

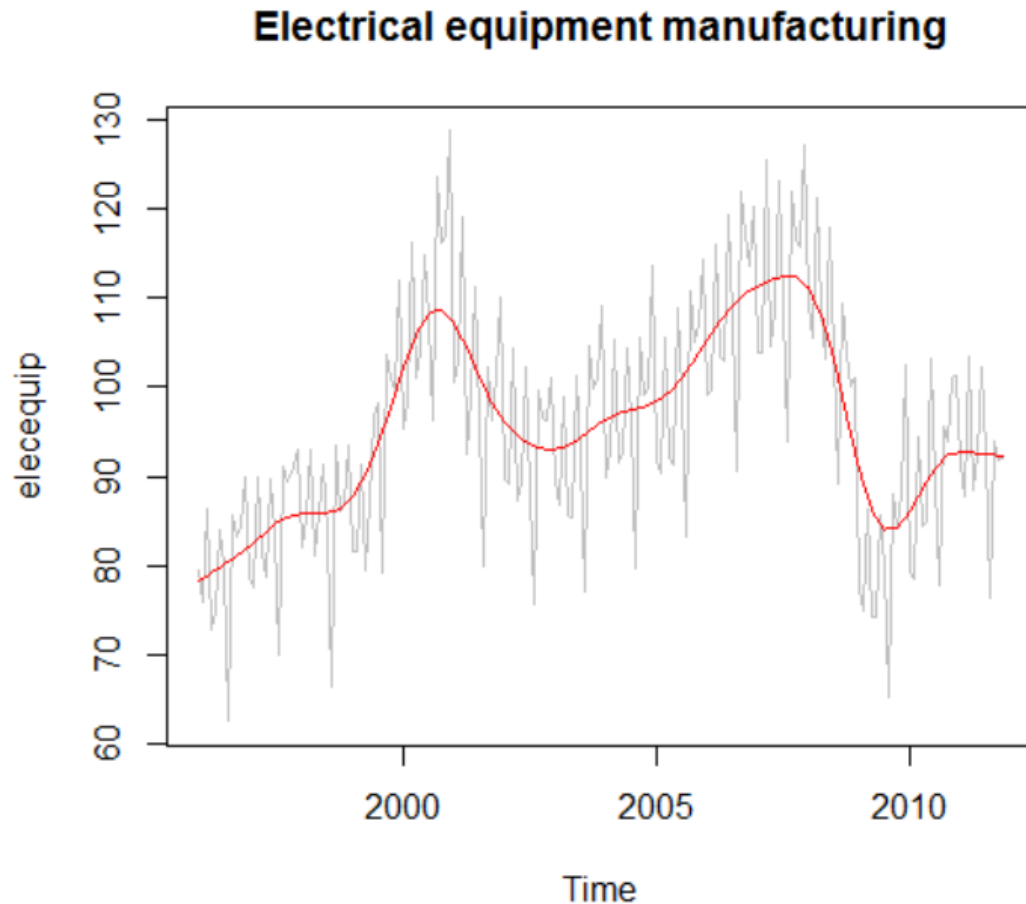
# Example: Seasonally-adjusted TS

```
fit <- stl ( elecequip, s.window=5 )  
plot(elecequip, col="grey", main="Electrical equipment manufacturing")  
lines(seasadj (fit), col="red", ylab="Seasonally adjusted")
```



# Example: Trend/Cycle Component in TS

```
fit <- stl ( elecequip, s.window=5 )  
plot(elecequip, col="grey", main="Electrical equipment manufacturing")  
lines(fit$time.series[ , 2], col="red", ylab="Trend")
```



# Global vs Local Patterns in TS Data

---

- **Global patterns** are the patterns in the TS data that extend throughout the TS period used for training and forecasting:
  - Trends: persistent upward or persistent downward trend with a non-changing slope
- **Local patterns** change over-time in a hard-to-predict manner

# How to Detect, Extract & Quantify TS Parts?

---

- The answer to this question depends a number of assumptions:
  - Whether the forecasting model to learn is persistent, i.e., does not / will not change over time
  - Whether TS exhibits long temporal scale patterns (i.e., **global patterns**) or short temporal scale patterns (i.e., **local patterns**)
  - How frequent do the short-scale patterns change over time?



# TS Data Analysis Methods

## TS Data Analysis & Forecasting

```
graph TD; A[TS Data Analysis & Forecasting] --> B[Data-Driven]; A --> C[Model-based]
```

### Data-Driven

Data-driven methods are used when model assumptions are likely to be violated, or when the structure of time series changes over time.

- **Baseline:** average, naive, seasonal naive, drift
- **Differencing**
- **Smoothing:** moving average, exponential smoothing

### Model-based

Training data is used to estimate model parameters, and then the model with these parameters is used to generate forecasts.

- **ARIMA**
- **Linear Regression**
- **Logistic Regression**
- **Neural Networks**

# Data-driven vs. Model-Based Methods

- **Model-based methods** are generally preferable for forecasting series with global patterns; they use all the data to estimate the **global patterns**.
  - For a local pattern, a model would require specifying how and when the patterns change, which is usually impractical and often unknown.
  - Model-based methods such as neural networks, regression trees, etc. are also used for TS forecasting for **incorporating external information into forecasts**.
- **Data-driven methods** are preferable for forecasting series with **local patterns**. Such methods “learn” patterns from the data, and their memory length can be set to best adapt to the rate of change in the series.
  - Patterns that change quickly warrant a “short memory,” whereas patterns that change slowly warrant a “long memory”

# TS Data Analysis Methods

## TS Data Analysis & Forecasting

```
graph TD; A[TS Data Analysis & Forecasting] --> B[Data-Driven]; A --> C[Model-based];
```

### Data-Driven

Data-driven methods are used when model assumptions are likely to be violated, or when the structure of time series changes over time.

- **Baseline:** average, naive, seasonal naive, drift
- **Differencing**
- **Smoothing:** moving average, exponential smoothing

### Model-based

Training data is used to estimate model parameters, and then the model with these parameters is used to generate forecasts.

- **ARIMA**
- **Linear Regression**
- **Logistic Regression**
- **Neural Networks**

# Baseline: Simple Forecasting Methods

- **Average:** `meanf ( ts.data, h=20 )`
  - Forecast of all future values is the mean of historical data  $\{y_1, \dots, y_T\}$
  - $F_{T+h} = \hat{y}_{T+h} = \bar{y} = (y_1 + \dots + y_T)/T$
- **Naive:** `naive (ts.data, h=20)` or `rwf (ts.data, h=20)`
  - Forecast is equal to the last observed value
  - $F_{T+h|T} = \hat{y}_{T+h|T} = y_T$
- **Seasonal naive:** `snaive (ts.data, h=20)`
  - Forecast is equal to the last value from the same season
  - $\hat{y}_{T+h|T} = y_{T+h-km}$ , where  $m$  is the seasonal period and  $k = \text{round}\left(\frac{h-1}{m}\right) + 1$
- **Drift:** `rwf (ts.data, drift=TRUE, h=20)`
  - Forecast is equal to the last value plus the average change
  - Equivalent to extrapolating a line between the first and last observation
  - $F_{T+h|T} = \hat{y}_{T+h|T} = y_T + \frac{h}{T-1} \sum_{t=2}^T (y_t - y_{t-1}) = y_T + \frac{h}{T-1} (y_T - y_1)$

---

# Removing Trend & Seasonality



## **DIFFERENCING**

# Differencing

- **Differencing:** take the difference between two time series.
  - Used for removing a trend and/or a seasonal pattern
  - **De-trend:** to remove the trend in a ts
  - **De-seasonalize:** to remove seasonality in a ts
- **De-trend:**
  - A **lag-1** difference ( $y_t - y_{t-1}$ ) means taking the difference between every two consecutive values in a series.
  - lag-1 differencing results in series that measure the changes from one period to the next
- **De-seasonalize:**
  - Differencing at lag- $k$  means subtracting the value from  $k$  periods back ( $y_t - y_{t-k}$ )
  - Example: lag-12 difference of the Amtrak ridership to remove annual seasonality: `diff(ridership.ts, lag = 12)`
- **Double Differencing:**
  - To remove both trend and seasonality
  - To remove higher-order trends



# Differencing for de-trending

- lag-1 differencing is useful for removing a trend
- De-trending with differencing vs. regression:
  - differencing does not assume that the trend is global, i.e., that the trend is fixed throughout the entire period.
- For *quadratic* and *exponential* trends, often another round of lag-1 differencing must be applied in order to remove the trend



# Differencing for Removing Seasonality


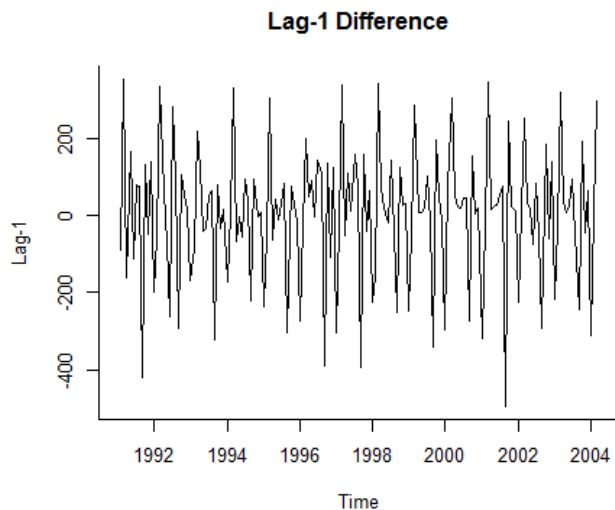


- To remove a seasonal pattern with  $M$  seasons, difference at lag  $M$ .
  - to remove a day-of-week pattern in daily data, take lag-7 differences
- Example, lag-12 applied to Amtrak monthly data removes monthly pattern: `diff(ridership.ts, lag = 12)`






# Removing Trend and Seasonality

- When both trend and seasonality, apply differencing twice



# Smoothing Methods

## **MOVING AVERAGE**



# Smoothing Methods

---

- Smoothing methods “smooth” out the noise in a series in an attempt to uncover the patterns.
- Smoothing is done by averaging over multiple observations
- **Different smoothers differ by**
  - the number of observations averaged,
  - how the average is computed,
  - how many times averaging is performed

# Moving-Average Smoothing


- Moving average consists of averaging across a window of consecutive observations generating a series of averages
- Two types of moving averages:
  - a **centered** moving average
  - a **trailing** moving average
- **Centered** moving averages are used for **visualizing trends**:
  - the averaging suppresses seasonality and noise, thus
  - making the trend more visible
- **Trailing** moving averages are used for **forecasting**

**Use Trailing Moving-Average for **forecasting** ONLY after both **trends and seasonality** are removed.**

# Centered Moving-Average for Visualization


- In a **centered moving-average** the value of the **moving average at time  $t$**  ( $MA_t$ ) is computed by centering the window around time  $t$  and averaging across the  $w$  values within the window

$$MA_t = \frac{y_{t-\frac{w-1}{2}} + \cdots y_{t-1} + y_t + y_{t+1} + \cdots + y_{t+\frac{w-1}{2}}}{w}$$



# Window Width Selection


- To choose the window width in a seasonal series:
  - because the goal is to suppress seasonality for better visualizing the trend, the default choice is **the length of a seasonal cycle**
- Example:
  - Amtrak ridership data, the annual seasonality indicates  $w = 12$



# Trailing Moving-Average for Forecasting

- Centered moving averages are computed by averaging across data both in the past and future of a given time point.
  - They cannot be used for forecasting because at the time of forecasting, the future is typically unknown
- For the purposes of forecasting, use **trailing moving averages**:
  - where the window of width  $w$  is set on the most recent available  $w$  of the series.
- The  $k$ -step-ahead forecast  $F_{t+k}$  ( $k = 1, 2, 3..$ ) is then the average of the these  $w$  values

$$F_{t+k} = \frac{y_t + y_{t-1} + y_{t-2} + \cdots + y_{t-w+1}}{w}$$



# Example: Moving Average Smoothing

```
library("forecast")
library("zoo")
Amtrak.data <- read.csv("Amtrak data.csv")
ridership.ts <- ts(Amtrak.data$Ridership, start = c(1991, 1),
                  end = c(2004, 3), freq = 12)

ma.trailing <- rollmean(ridership.ts, k = 12, align = "right")
ma.centered <- ma(ridership.ts, order = 12)
plot(ridership.ts, ylim = c(1300, 2200), ylab = "Ridership",
     xlab = "Time", bty = "l", xaxt = "n",
     xlim = c(1991, 2004.25), main = "")
axis(1, at = seq(1991, 2004.25, 1),
     labels = format(seq(1991, 2004.25, 1)))
lines(ma.centered, lwd = 2)
lines(ma.trailing, lwd = 2, lty = 2)
legend(1994, 2200, c("Ridership", "Centered Moving Average",
                    "Trailing Moving Average"),
     lty=c(1,1,2), lwd=c(1,2,2), bty = "n")
```



# Notes: Lagging and De-trending

- Moving average “lags behind,” thereby it is **over-forecasting** in the presence of an increasing trend and **under-forecasting** in the presence of a decreasing trend:
  - In general, the moving average can be used for forecasting only in a series that lacks seasonality and trend.
- To remove trends (de-trending) and to remove seasonality (de-seasonalizing) from a time series, use the following methods:
  - regression models
  - advanced exponential smoothing methods
  - differencing
- Use the moving average to forecast such de-trended and de-seasonalized series, and
- Finally, add the trend and seasonality back to the forecast

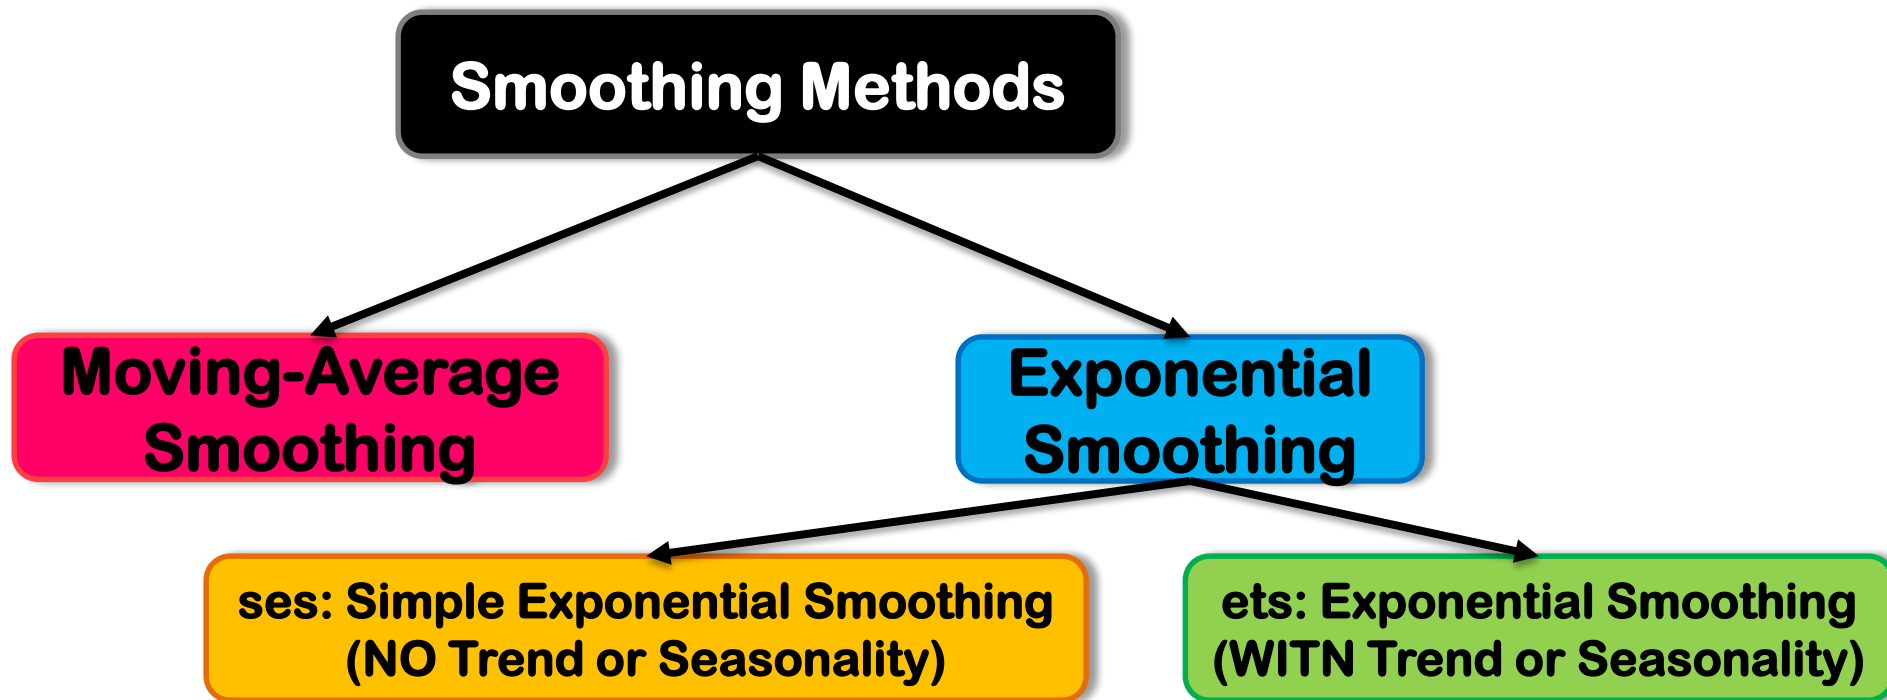
# Steps with Moving-Average Forecasting

---

1. De-trend
2. De-seasonalize
3. Compute  $F_{t+k}$  trailing moving-average forecast
4. Add trend and seasonality back to produce the forecast,  $\widehat{y}_{t+k}$ :

# Smoothing Methods

## EXPONENTIAL SMOOTHING



# Simple Exponential Smoothing for Forecasting

- **Simple exponential smoothing (SES)** is a **weighted average of all past values**, so that weights decrease exponentially into the past:
  - unlike forecasting with a moving average of the  $w$  most recent values
  - it weights more recent values more heavily
- **Forecast with simple exponential smoothing only after removing trends or seasonality:**
  - i.e.: apply exponential smoothing **to the series of residuals**
- **Forecast at time  $t + 1$  ( $F_{t+1}$ ) with exponential smother as follows:**

$$F_t = \alpha y_t + \alpha(1 - \alpha)y_{t-1} + \alpha(1 - \alpha)^2 y_{t-2} + \dots$$

$$F_t = \alpha y_t + (1 - \alpha)F_{t-1}$$

recursive definition

where  $\alpha$  is a constant between 0 and 1 called **smoothing constant**

**Use Exponential Smoothing for forecasting ONLY after both trends and seasonality are removed.**

# Exponential Forecaster: Active Learner

- The exponential smoother is a weighted average of all past observations, with exponentially decaying weights.
- Practically it is better to write it in another way:

$$F_{t+1} = F_t + \alpha e_t$$

- where  $e_t$  is the forecast error at time  $t$
- This formulation presents the exponential forecaster as an “active learner”
  - It looks at the previous forecast  $F_t$  and its distance from the actual value  $e_t$  and then corrects the next forecast based on that information.
    - If the forecast is too high in the last period, the next period is adjusted down.
    - The amount of correction depends on the value of the smoothing constant  $\alpha$

# Active Learning Formulation

- **Formulation**

$$F_{t+1} = F_t + \alpha e_t$$

- **is advantageous in terms of data storage and computation time**
  - need to store only forecast and forecast error from the previous period, not entire series
  - For real-time forecasting, or forecasting many series in parallel and continuously, such savings are critical
- **The forecasting further into the future yields the same forecasts as a one-step-ahead forecast.**
  - Because series is assumed to lack trend and seasonality, forecasts into the future rely only on information that is available at the time of prediction

$$F_{t+k} = F_{t+1}$$

# Choosing Smoothing Constant, $\alpha$

- The smoothing constant  $\alpha$ , which is set by the user, determines the rate of learning.
  - A value close to 1 indicates fast learning
    - only the most recent observations have influence on forecasts
  - A value close to 0 indicates slow learning
    - past observations have a large influence on forecasts
  - Default values: in the range of 0.1-0.2.
- Trial and error can also help in the choice of  $\alpha$ 
  - Examine the time plot of the actual and predicted series, as well as predictive accuracy
  - MAPE or RMSE of the validation period
- Finding the  $\alpha$  value that optimizes one-step-ahead predictive accuracy over the training period can be used to determine the degree of *local vs. global nature* of the level.
  - beware of choosing the “**best  $\alpha$** ” for forecasting as this can lead to model **overfitting** and low predictive accuracy over in validation

# ets() for Simple Exponential Smoothing in R


---

- In R, forecasting using simple exponential smoothing can be done via the `ets()` function
- `ets` stands for **E**rror, **T**rend, and **S**easonality
- The **simple exponential smoothing** model=“**ANN**”
  - additive error (A)
  - no trend (N)
  - no seasonality (N)



# Example: Simple Exponential Smoothing

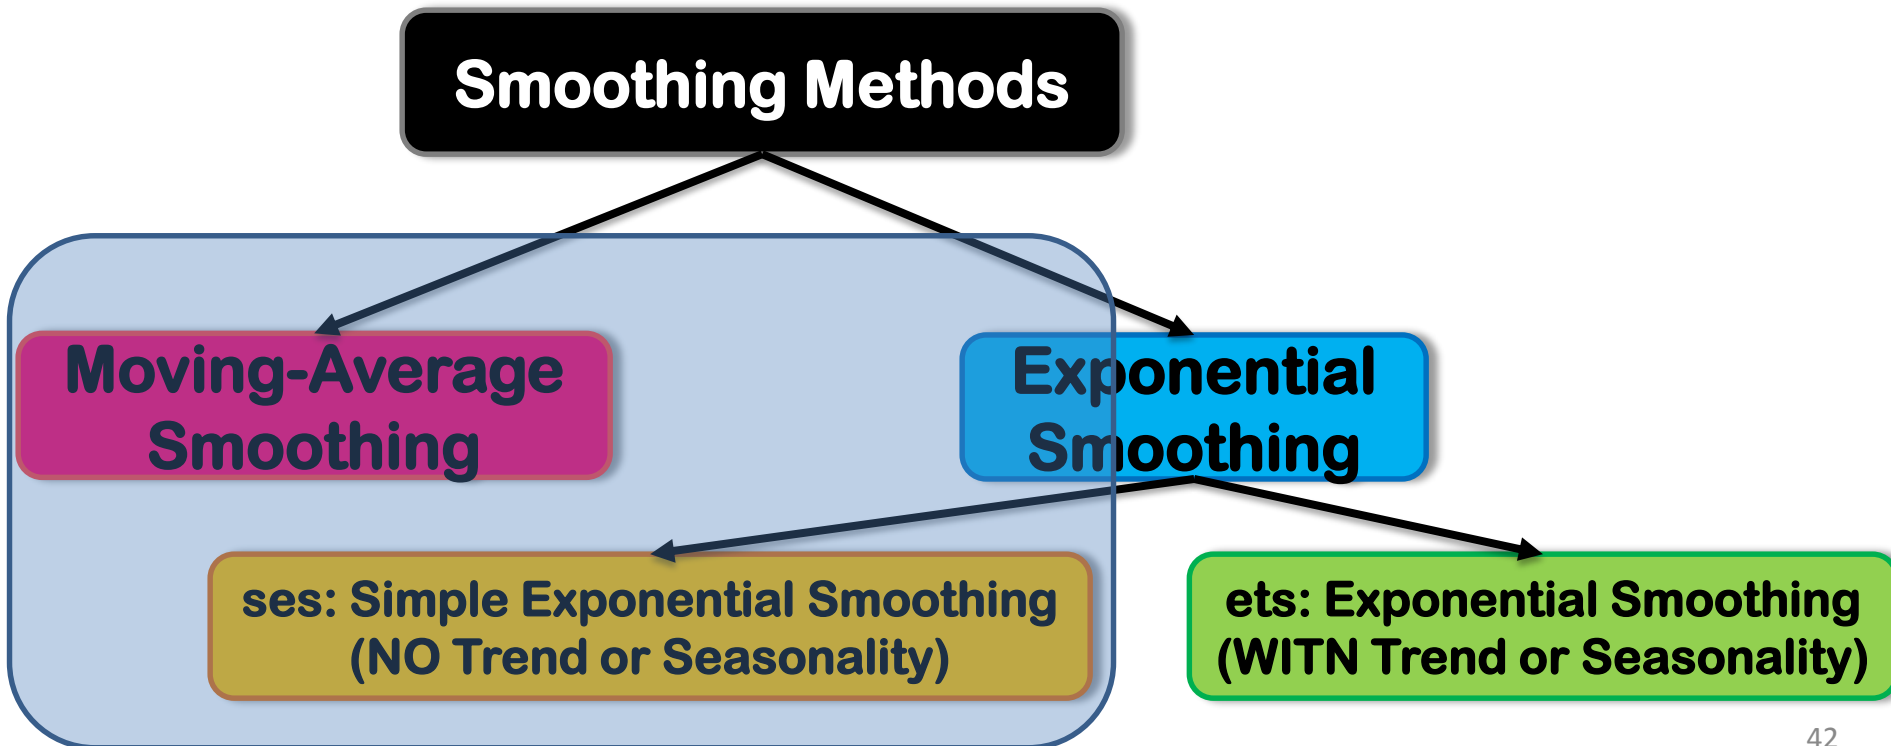
```
Amtrak.data <- read.csv("Amtrak data.csv")
ridership.ts <- ts(Amtrak.data$Ridership,
  start = c(1991, 1),
  end = c(2004, 3), freq = 12)
```



```
diff.twice.ts <- diff(diff(ridership.ts, lag = 12), lag = 1)
nValid <- 36
nTrain <- length(diff.twice.ts) - nValid
train.ts <- window(diff.twice.ts, start = c(1992, 2),
  end = c(1992, nTrain + 1))
valid.ts <- window(diff.twice.ts, start = c(1992, nTrain + 2),
  end = c(1992, nTrain + 1 + nValid))
ses <- ets(train.ts, model = "ANN", alpha = 0.2)
ses.pred <- forecast(ses, h = nValid, level = 0)
```

# Smoothing Methods

## EXPONENTIAL SMOOTHING



# Advanced Exponential Smoothing

---

- Both moving average and simple exponential smoothing should only be used for forecasting series with no trend or seasonality
- One approach is to remove trend and seasonality before smoothing
- Another approach is **to use more sophisticated version of exponential smoothing that captures trend and/or seasonality**

# Holt's Model: Series with an Additive Trend

- **Holt's linear trend model**
  - For series that contain **a trend**
  - Uses double exponential smoothing
  - Trend is not assumed to be global, but can change over time
- **Double Exponential Smoothing:**
  - the local trend is estimated from the data, and is updated as more data become available
  - the level of the series is also estimated from the data, and is updated as more data become available
- **The k-step-ahead forecast is given by combining the level estimate at time  $t$  ( $L_t$ ) and the trend estimate (which is assumed to be additive) at time  $t$  ( $T_t$ )**

$$F_{t+k} = L_t + kT_t$$

# k-step-ahead Forecast

$$F_{t+k} = L_t + kT_t$$

- In the presence of a trend, one- two- three-step-ahead forecasts are no longer identical.
- Update the level and trend as follows:

$$L_t = \alpha y_t + (1 - \alpha)(L_{t-1} + T_{t-1})$$

$$T_t = \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}$$

- The level at time  $t$  is a weighted average of the actual value at time  $t$  and the level in the previous period, adjusted for the trend
  - The trend at time  $t$  is a weighted average of the trend in the previous period and the more recent information on the change in level.
- Two smoothing constants  $\alpha$  and  $\beta$  determine the rate of learning

# Two Types of Errors in Exponential Smoothing

- **Additive error**

- $y_{t+1}$  consists not only of level+trend but also an additional error:
  - $y_{t+1} = L_t + T_t + e_t$
- the errors are assumed to have a fixed magnitude, irrespective of the current level+trend of the series

- **Multiplicative error**

- The size of the error grows as the level of the series increases
  - $y_{t+1} = (L_t + T_t) \times (1 + e_t)$
- Error acts as a percentage increase in the current level+trend

- In R, `etc()` function includes either option:

- `model = "AAN"` or
- `model = "MAN"`
- The first letter represents **error type**

# Holt-Winter's Model: Trend + Seasonality

- The **Holt-Winter's exponential smoothing method**:
  - For series that contain **both trend and seasonality**
- **In multiplicative seasonality**:
  - values on different seasons differ by percentage amounts
- **In additive seasonality**:
  - values on different seasons differ by a fixed amount
- **Assuming seasonality with  $M$  seasons, the forecast for an additive trend and multiplicative seasonality is given by:**

$$F_{t+k} = (L_t + kT_k)S_{t+k-M}$$

- By forecasting time  $t$ , the series must include at least a full cycle of seasons in order to produce forecasts using this formula ( $t > M$ )

# Update in Holt-Winter's Method

- Being an adaptive method, Holt-Winter's exponential smoothing allows the level, trend and seasonality patterns to change over time
- The three components are estimated and updated as new information arrives
- The updating equations for this additive trend and multiplicative seasonality are:

$$L_t = \frac{\alpha y_t}{S_{t-M}} + (1 - \alpha)(L_{t-1} + T_{t-1})$$

$$T_t = \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}$$

$$S_t = \gamma \left( \frac{y_t}{L_t} \right) + (1 - \gamma)S_{t-M}$$



# Summary: Exponential Smoothing Methods

Method	Trend	Season	Example
Simple exponential smoothing	NO	NO	ses(x)
Holt's method	Linear	NO	holt(x)
Exponential trend	Exponential	NO	holt(x, exponential=TRUE)
Damped trend	Linear	NO	holt(x, damped=TRUE)
Damped exponential trend	Exponential	NO	holt(x, damped=TRUE, exponential=TRUE)
Holt and Winter's	YES	YES	hw(x)

# Exponential Smoothing with `ets()` in R

- 2 error types × 3 trend types × 3 seasonality types
- Automated model selection
  - by leaving out the model option altogether or by using `model="ZZZ"`, the `ets()` function will fit several models and choose the “best” one
  - You can also automatically select a model from a subset of models by putting Z in one or two places in the model specification
    - e.g. `model = "AZZ"` automatically looks for best **additive error** model
- “Best” is by **minimizing Akaike’s Information Criterion (AIC)**
  - combines fit to the training data with a penalty for the number of smoothing parameters and initial values ( $L_0, T_0, etc$ ) included in the model
- Selection Metrics:
  - **$AIC = -\log(likelihood) + 2p$  (Akaike’s Information Criterion)**
    - where likelihood is a measure of fit to the training period and  $p$  is the number of smoothing parameters and initial states in the model
  - **$BIC = AIC + p(\log(n) - 2)$  (Bayesian Information Criterion)**

# Taxonomy of Exponential Smoothing Methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

(N,N): Simple exponential smoothing

(A,N): Holt's linear method

(A,A): Additive Holt-Winters' method

(A,M): Multiplicative Holt-Winters' method

(A<sub>d</sub>,M): Damped multiplicative Holt-Winters' method

# Taxonomy of Exponential Smoothing Methods

		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

R functions	Methods
<b>ses()</b>	(N,N)
<b>holt()</b>	(A,N), (A <sub>d</sub> ,N), (M,N), (M <sub>d</sub> ,N)
<b>hw()</b>	(A,A), (A <sub>d</sub> ,A), (A,M), (A <sub>d</sub> ,M), (M,M), (M <sub>d</sub> ,M)

# Automated model selection in ets()

- Here AIC difference of 9.139 strongly favors the MNA model

```
> ets(train.ts)
ETS(M,N,A)

call:
ets(y = train.ts)

Smoothing parameters:
  alpha = 0.5653
  gamma = 1e-04

Initial states:
  l = 1899.2864
  s=25.8874 -10.448 -0.6701 -122.7463 199.0173 150.1639
      42.8476 82.9771 52.7806 48.1995 -254.7762 -213.2329

sigma: 0.0316

      AIC      AICc      BIC
1608.976 1612.865 1648.347
> |
```

# AIC/BIC

- **Selection Metrics:**

- $AIC = -\log(\text{likelihood}) + 2p$  (**Akaike's Information Criterion**)
  - where likelihood is a measure of fit to the training period and  $p$  is the number of smoothing parameters and initial states in the model
- $BIC = AIC + p(\log(n) - 2)$  (**Bayesian Information Criterion**)

- **How to interpret AIC/BIC:**

- AIC/AICs/BIC does not have much meaning by itself.
- Only useful in comparison to AIC value for another model fitted to *same data set*.
- Consider several models with AIC values close to the minimum.
- A difference in AIC values of 2 or less is not regarded as substantial → choose the simpler but non-optimal model.
- AIC can be negative.

# Example: Using ets() Best Model Selection

```
fit <- ets(ausbeer)
fit2 <- ets(ausbeer,model="AAA",damped=FALSE)
fcast1 <- forecast(fit, h=20)
fcast2 <- forecast(fit2, h=20)
```

```
ets(y, model="ZZZ", damped=NULL, alpha=NULL,
    beta=NULL, gamma=NULL, phi=NULL,
    additive.only=FALSE,
    lower=c(rep(0.0001,3),0.80),
    upper=c(rep(0.9999,3),0.98),
    opt.crit=c("lik","amse","mse","sigma"), nmse=3,
    bounds=c("both","usual","admissible"),
    ic=c("aic","aicc","bic"), restrict=TRUE)
```

# Example: ets() (cont.)

```
> fit  
ETS(M,Md,M)
```

Smoothing parameters:

alpha = 0.1776

beta = 0.0454

gamma = 0.1947

phi = 0.9549

Initial states:

l = 263.8531

b = 0.9997

s = 1.1856 0.9109 0.8612 1.0423

sigma: 0.0356

AIC	AICc	BIC
2272.549	2273.444	2302.715

```
> fit2  
ETS(A,A,A)
```

Smoothing parameters:

alpha = 0.2079

beta = 0.0304

gamma = 0.2483

Initial states:

l = 255.6559

b = 0.5687

s = 52.3841 -27.1061 -37.6758 12.3978

sigma: 15.9053

AIC	AICc	BIC
2312.768	2313.481	2339.583



# Summary of ets() function


---

**ETS**: **E**rror, **T**rend, **S**easonal

- **Automatically chooses** a model by default using the AIC/AICc/BIC
- Can handle any **combination of trend, seasonality and damping**
- Produces prediction intervals for every model
- Ensures the parameters are admissible
- In practice, the models work fine for short- to medium-term forecasts provided the data are strictly positive.
- Produces an object of class ets:
  - Methods: coef(), plot(), summary(), residuals(), fitted(), simulate(), and forecast()
  - plot() function shows time plots of the original time series along with the extracted components (level, growth and seasonal).

# plot(fit)

## Decomposition by ETS(M,Md,M) method



# Accuracy (more in the future lecture)

---


```
> accuracy(fit)
```

ME	RMSE	MAE	MPE	MAPE	MASE
0.17847	15.48781	11.77800	0.07204	2.81921	0.20705

```
> accuracy(fit2)
```

ME	RMSE	MAE	MPE	MAPE	MASE
-0.11711	15.90526	12.18930	-0.03765	2.91255	0.21428

# Forecast from ETS()



# Acknowledgements

---

- **Books**

- Free and online ([otexts.com/fpp](http://otexts.com/fpp)): Forecasting Principles & Practice by R. Hyndman, G. Athanasopoulos ← **Excellent Book!!!**
- Practical Time Series Forecasting with R: A Hand-on Guide by Shmueli & Lichtendahl

- **Packages**

- R: fpp (`install.packages("fpp", dependencies=TRUE)`)