


Chapter 18: Concurrency Control

(Slides by Hector Garcia-Molina,
<http://www-db.stanford.edu/~hector/cs245/notes.htm>)

Chapter 18

1

Chapter 18 Concurrency Control



Chapter 18

2

Example:

T1: Read(A)	T2: Read(A)
$A \leftarrow A + 100$	$A \leftarrow A \times 2$
Write(A)	Write(A)
Read(B)	Read(B)
$B \leftarrow B + 100$	$B \leftarrow B \times 2$
Write(B)	Write(B)
Constraint: $A = B$	

Chapter 18

3

Schedule A		A	B
T1	T2	25	25
Read(A); A \leftarrow A+100			
Write(A);		125	
Read(B); B \leftarrow B+100;			125
Write(B);	Read(A); A \leftarrow A \times 2;		
	Write(A);	250	
	Read(B); B \leftarrow B \times 2;		250
	Write(B);	250	250

Chapter 18

4

Schedule B		A	B
T1	T2	25	25
	Read(A); A \leftarrow A \times 2;		
	Write(A);	50	
	Read(B); B \leftarrow B \times 2;		50
	Write(B);		
Read(A); A \leftarrow A+100		150	
Write(A);			150
Read(B); B \leftarrow B+100;			
Write(B);		150	150

Chapter 18

5

Schedule C		A	B
T1	T2	25	25
Read(A); A \leftarrow A+100			
Write(A);		125	
	Read(A); A \leftarrow A \times 2;		
	Write(A);	250	
Read(B); B \leftarrow B+100;			125
Write(B);	Read(B); B \leftarrow B \times 2;		250
	Write(B);	250	250

Chapter 18

6

Schedule D

T1	T2	A	B
Read(A); A \leftarrow A+100		25	25
Write(A);		125	
	Read(A); A \leftarrow A \times 2;	250	
	Write(A);		
	Read(B); B \leftarrow B \times 2;		50
	Write(B);		150
Read(B); B \leftarrow B+100;		250	150
Write(B);			

Chapter 18

7

Schedule E

Same as Schedule D
but with new T2'

T1	T2'	A	B
Read(A); A \leftarrow A+100		25	25
Write(A);		125	
	Read(A); A \leftarrow A \times 1;	125	
	Write(A);		
	Read(B); B \leftarrow B \times 1;		25
	Write(B);		125
Read(B); B \leftarrow B+100;		125	125
Write(B);			

Chapter 18

8

- Want schedules that are "good", regardless of
 - initial state and
 - transaction semantics
- Only look at order of read and writes

Example:

Sc=r₁(A)w₁(A)r₂(A)w₂(A)r₁(B)w₁(B)r₂(B)w₂(B)

Chapter 18

9

Example:

$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$

$Sc' = r_1(A)w_1(A) \underbrace{r_1(B)w_1(B)r_2(A)w_2(A)}_{T_1} \underbrace{r_2(B)w_2(B)}_{T_2}$

Chapter 18

10

However, for Sd :

$Sd = r_1(A)w_1(A)r_2(A)w_2(A) \underbrace{r_2(B)w_2(B)r_1(B)w_1(B)}_{\text{X}}$

Chapter 18

11

- $T_2 \rightarrow T_1$
- Also, $T_1 \rightarrow T_2$

$T_1 \rightarrow T_2$ \Rightarrow Sd cannot be rearranged into a serial schedule
 \Rightarrow Sd is not "equivalent" to any serial schedule
 \Rightarrow Sd is "bad"

Chapter 18

12

Returning to Sc

$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$

$\swarrow \quad \searrow \quad \quad \quad \swarrow \quad \searrow$

$T_1 \rightarrow T_2 \quad \quad \quad T_1 \rightarrow T_2$

- no cycles \Rightarrow Sc is "equivalent" to a serial schedule (in this case T_1, T_2)

Chapter 18

13

Concepts

Transaction: sequence of $r_i(x)$, $w_i(x)$ actions

Conflicting actions:

$\begin{matrix} r_1(A) & & w_2(A) & & w_1(A) \\ & \swarrow \quad \searrow & & \swarrow \quad \searrow & \\ & w_2(A) & r_1(A) & w_2(A) & \end{matrix}$

Schedule: represents chronological order in which actions are executed

Serial schedule: no interleaving of actions or transactions

Chapter 18

14

Definition

S_1, S_2 are conflict equivalent schedules if S_1 can be transformed into S_2 by a series of swaps on non-conflicting actions.

Chapter 18

15

Definition

A schedule is conflict serializable if it is conflict equivalent to some serial schedule.

Chapter 18

16

Precedence graph $P(S)$ (S is schedule)

Nodes: transactions in S

Arcs: $T_i \rightarrow T_j$ whenever

- $p_i(A), q_j(A)$ are actions in S
- $p_i(A) <_S q_j(A)$
- at least one of p_i, q_j is a write

Chapter 18

17

Exercise:

- What is $P(S)$ for
 $S = w_3(A) w_2(C) r_1(A) w_1(B) r_1(C) w_2(A) r_4(A) w_4(D)$

- Is S serializable?

Chapter 18

18

Lemma

S_1, S_2 conflict equivalent $\Rightarrow P(S_1)=P(S_2)$

Chapter 18

19

Note: $P(S_1)=P(S_2) \not\Rightarrow S_1, S_2$ conflict equivalent

Counter example:

$S_1 = w_1(A) \ r_2(A) \quad w_2(B) \ r_1(B)$

$S_2 = r_2(A) \ w_1(A) \quad r_1(B) \ w_2(B)$

Chapter 18

20

Theorem

$P(S_1)$ acyclic $\Leftrightarrow S_1$ conflict serializable

Chapter 18

21

How to enforce serializable schedules?


Option 1: run system, recording P(S);
at end of day, check for P(S)
cycles and declare if execution
was good

Chapter 18

22

How to enforce serializable schedules?

Option 2: prevent P(S) cycles from
occurring



Chapter 18


23

A locking protocol

Two new actions:

lock (exclusive): $li(A)$

unlock: $ui(A)$



Chapter 18

24

Rule #1: Consistent transactions

$T_i: \dots l_i(A) \dots p_i(A) \dots u_i(A) \dots$

Chapter 18

25

Rule #2 Legal scheduler

$S = \dots l_i(A) \dots u_i(A) \dots$
 $\quad \quad \quad \longleftrightarrow$
 $\quad \quad \quad \text{no } l_j(A)$

Chapter 18

26

Exercise:

- What schedules are legal?
What transactions are consistent?
 $S_1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$
 $S_2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$
 $S_3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

Chapter 18

27

Schedule F

T1	T2
l ₁ (A);Read(A)	
A ← A + 100;Write(A);u ₁ (A)	
	l ₂ (A);Read(A)
	A ← A × 2;Write(A);u ₂ (A)
	l ₂ (B);Read(B)
	B ← B × 2;Write(B);u ₂ (B)
l ₁ (B);Read(B)	
B ← B + 100;Write(B);u ₁ (B)	

Chapter 18

28

Schedule F

T1	T2	A	B
l ₁ (A);Read(A)		25	25
A ← A + 100;Write(A);u ₁ (A)		125	
	l ₂ (A);Read(A)	250	
	A ← A × 2;Write(A);u ₂ (A)		
	l ₂ (B);Read(B)		50
	B ← B × 2;Write(B);u ₂ (B)		150
l ₁ (B);Read(B)		250	150
B ← B + 100;Write(B);u ₁ (B)			

Chapter 18

29


Rule #3 Two phase locking (2PL) for transactions

T_i = l_i(A) u_i(A)



Chapter 18

30



Schedule G

T1	T2
l ₁ (A); Read(A)	
A ← A + 100; Write(A)	
l ₁ (B); u ₁ (A)	
	l ₂ (A); Read(A)
	A ← A x 2; Write(A); l ₂ (B)

Chapter 18

32

Schedule G

T1	T2
l ₁ (A); Read(A)	
A ← A + 100; Write(A)	
l ₁ (B); u ₁ (A)	
	l ₂ (A); Read(A)
	A ← A x 2; Write(A); l ₂ (B)
Read(B); B ← B + 100	
Write(B); u ₁ (B)	

Chapter 18

33

Schedule G

T1	T2
$I_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$I_1(B); u_1(A)$	
	$I_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); I_2(B)$
$\text{Read}(B); B \leftarrow B + 100$	
$\text{Write}(B); u_1(B)$	
	$I_2(B); u_2(A); \text{Read}(B)$
	$B \leftarrow B \times 2; \text{Write}(B); u_2(B);$

Chapter 18

34

Schedule H (T₂ reversed)

T1	T2
$I_1(A); \text{Read}(A)$	$I_2(B); \text{Read}(B)$
$A \leftarrow A + 100; \text{Write}(A)$	$B \leftarrow B \times 2; \text{Write}(B)$
$I_1(B)$	$I_2(A)$
delayed	delayed

deadlock

Chapter 18

35

- Assume deadlocked transactions are rolled back
 - They have no effect
 - They do not appear in schedule

E.g., Schedule H on previous slide

Chapter 18

36

We can show that
rules #1,2,3 \Rightarrow conflict-serializable
schedules

(see textbook)

Chapter 18

37

- Beyond this simple 2PL protocol, it is all a matter of improving performance and allowing more concurrency....
 - Shared locks
 - Multiple granularity
 - Inserts, deletes and phantoms
 - Other types of C.C. mechanisms

Chapter 18

38

Shared locks

So far:

$S = \dots l_1(A) \ r_1(A) \ u_1(A) \ \dots \ l_2(A) \ r_2(A) \ u_2(A) \ \dots$

Do not conflict

Instead:

$S = \dots ls_1(A) \ r_1(A) \ ls_2(A) \ r_2(A) \ \dots \ us_1(A) \ us_2(A)$

Chapter 18

39

Lock actions

$l-t(A)$: lock A in t mode (t is S or X)

$u-t(A)$: unlock t mode (t is S or X)

Shorthand:

$u(A)$: unlock whatever modes

T_i has locked A

Chapter 18

40

- What about transactions that read and write same object?

Option 1: Request exclusive lock

$T_i = \dots l-X_1(A) \dots r_1(A) \dots w_1(A) \dots u_1(A) \dots$

Chapter 18

41

- What about transactions that read and write same object?

Option 2: Upgrade

(E.g., need to read, but don't know if will write...)

$T_i = \dots l-S_1(A) \dots r_1(A) \dots l-X_1(A) \dots w_1(A) \dots u_1(A) \dots$

The same as:
- Getting 2nd lock on A, or
- Dropping S, getting X lock

Chapter 18

42

Compatibility matrix

Comp

	S	X
S	true	false
X	false	false

Chapter 18

43

Lock types beyond S/X

Examples:

- (1) increment lock
(see textbook)
- (2) update lock

Chapter 18

44

Upgrade locks

Common deadlock problem with upgrades:

T1	T2
I-S ₁ (A)	
	I-S ₂ (A)
I-X ₁ (A)	
	I-X ₂ (A)
--- Deadlock ---	

Chapter 18

45

Solution

If T_i wants to read A and knows it may later want to write A, it requests update lock (not shared)

Chapter 18

46

How does locking work in practice?


- Every system is different
(E.g., may not even provide CONFLICT-SERIALIZABLE schedules)
- But here is one (simplified) way ...

Chapter 18

47

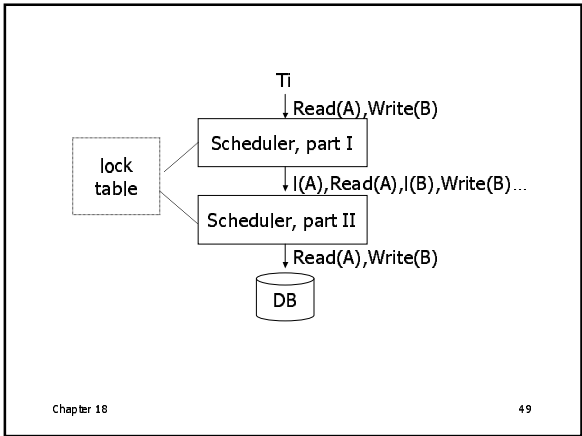
Sample Locking System:

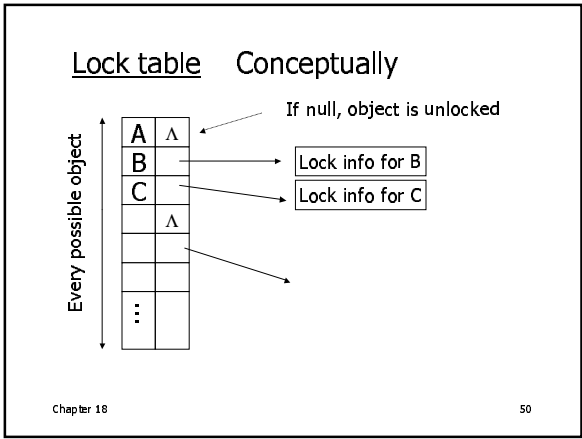
- (1) Don't trust transactions to request/release locks
- (2) Hold all locks until transaction commits

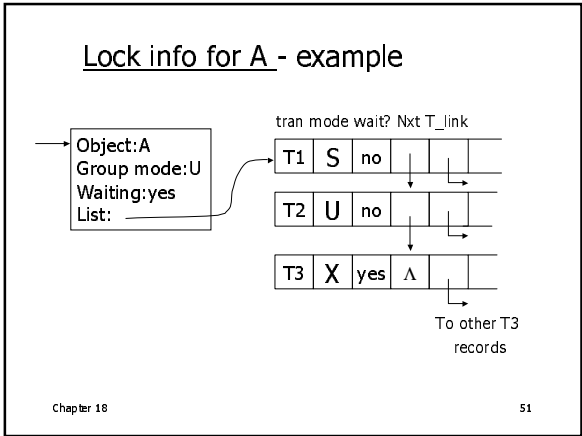


Chapter 18


48







What are the objects we lock?



Chapter 18

52

- Locking works in any case, but should we choose small or large objects?
- If we lock large objects (e.g., relations)
 - Need few locks
 - Low concurrency
- If we lock small objects (e.g., tuples, fields)
 - Need more locks
 - More concurrency

Chapter 18

53

Summary

Have looked at 2PL, a concurrency-control mechanism commonly used in practice

Others (in the textbook):

- Multiple granularity
- Tree (index) protocols
- Timestamping
- Validation

Chapter 18

54
