Relational Database Design: E/R-Relational Translation

Introduction to Databases
CompSci 316 Fall 2015

DUKE
COMPUTER SCIENCE

Announcements (Thu. Sep. 3)

- Homework #1 due on the 15th (the Tuesday after the next)
  - Please please please start early
- Project description available next week

Database design steps: review

- Understand the real-world domain being modeled
- Specify it using a database design model (e.g., E/R)
- Translate specification to the data model of DBMS (e.g., relational)
- Create DBMS schema

☞Next: translating E/R design to relational schema

E/R model: review

- Entity sets
  - Keys
  - Weak entity sets
- Relationship sets
  - Attributes on relationships
  - Multiplicity
  - Roles
  - Binary versus n-ary relationships
    - Modeling n-ary relationships with weak entity sets and binary relationships
 - ISA relationships

---

Translating entity sets

- An entity set translates directly to a table
  - Attributes → columns
 - Key attributes → key columns

User (*uid*, name)   Group (*gid*, name)

Translating weak entity sets

- Remember the "borrowed" key attributes
- Watch out for attribute name conflicts

Building (*name*, year)
Room (*building_name*, *room_number*, capacity)
Seat (*building_name*, *room_number*, *seat_number*, left_or_right)

Translating relationship sets

- A relationship set translates to a table
  - Keys of connected entity sets → columns
  - Attributes of the relationship set (if any) → columns
 - Multiplicity of the relationship set determines the key of the table

Member (uid, gid, fromDate)

More examples

Parent (parent_uid, child_uid)

*Member (uid, initiator_uid, gid)*

Translating double diamonds?

- Recall that a double-diamond (supporting) relationship set connects a weak entity set to another entity set
- No need to translate because the relationship is implicit in the weak entity set's translation


RoomInBuilding
(room_building_name, room_number, building_name)

is subsumed by

Room (building_name, room_number, capacity)

Translating subclasses & ISA: approach 1

- **Entity-in-all-superclasses** approach ("E/R style")
  - An entity is represented in the table for each subclass to which it belongs
  - A table includes only the attributes directly attached to the corresponding entity set, plus the inherited key


⟨142, Bart⟩ ∈ *Group (gid, name)*
⟨456, Ralph⟩ *User (uid, name)*
  Member (uid, gid, from_date)
⟨456, ☺⟩∈ *PaidUser (uid, avatar)*

Translating subclasses & ISA: approach 2

- **Entity-in-most-specific-class** approach ("OO style")
  - An entity is only represented in one table (the most specific entity set to which the entity belongs)
  - A table includes the attributes attached to the corresponding entity set, plus all inherited attributes

Group (gid, name)
⟨142, Bart⟩ ∈ *User (uid, name)*
  Member (uid, gid, from_date)
⟨456, Ralph, ☺⟩∈ *PaidUser (uid, name, avatar)*

Translating subclasses & ISA: approach 3
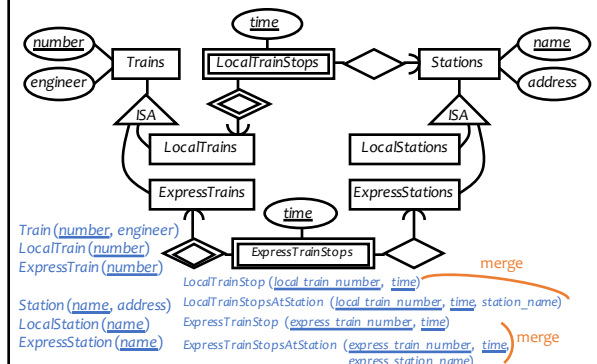
- **All-entities-in-one-table** approach ("NULL style")
  - One relation for the root entity set, with all attributes found in the network of subclasses (plus a "type" attribute when needed)
  - Use a special NULL value in columns that are not relevant for a particular entity

⟨142, Bart, NULL⟩ ∈ *Group (gid, name)*
⟨456, Ralph, ☺⟩ *User (uid, name, avatar)*
  Member (uid, gid, from_date)

Comparison of three approaches

- Entity-in-all-superclasses
  - *User* (*uid*, *name*), *PaidUser* (*uid*, *avatar*)
  - Pro: All users are found in one table
  - Con: Attributes of paid users are scattered in different tables
- Entity-in-most-specific-class
  - *User* (*uid*, *name*), *PaidUser* (*uid*, *name*, *avatar*)
  - Pro: All attributes of paid users are found in one table
  - Con: Users are scattered in different tables
- All-entities-in-one-table
  - *User* (*uid*, [*type*,] *name*, *avatar*)
  - Pro: Everything is in one table
  - Con: Lots of NULL's; complicated if class hierarchy is complex

A complete example

Train (*number*, *engineer*)
LocalTrain (*number*)
ExpressTrain (*number*)

Station (*name*, *address*)
LocalStation (*name*)
ExpressStation (*name*)

LocalTrainStop (*local_train_number*, *time*)
LocalTrainStopsAtStation (*local_train_number*, *time*, *station_name*)
ExpressTrainStop (*express_train_number*, *time*)
ExpressTrainStopsAtStation (*express_train_number*, *time*, *express_station_name*)

Simplifications and refinements

Train (*number*, *engineer*), *LocalTrain* (*number*), *ExpressTrain* (*number*)
Station (*name*, *address*), *LocalStation* (*name*), *ExpressStation* (*name*)
LocalTrainStop (*local_train_number*, *station_name*, *time*)
ExpressTrainStop (*express_train_number*, *express_station_name*, *time*)

- Eliminate *LocalTrain* table
  - Redundant: can be computed as
    $$\pi_{number}(Train) - ExpressTrain$$
  - Slightly harder to check that *local_train_number* is indeed a local train number
- Eliminate *LocalStation* table
  - It can be computed as $\pi_{number}(Station) - ExpressStation$

An alternative design

Train (*number*, *engineer*, *type*)

Station (*name*, *address*, *type*)

TrainStop (*train_number*, *station_name*, *time*)

- Encode the type of train/station as a column rather than creating subclasses
- What about the following constraints?
  - Type must be either "local" or "express"
  - Express trains only stop at express stations
  - ☞They can be expressed/declared explicitly as database constraints in SQL (as we will see later in course)
- Arguably a better design because it is simpler!

Design principles

POOR DESIGN!

- KISS
  - Keep It Simple, Stupid
- Avoid redundancy
  - Redundancy wastes space, complicates modifications, promotes inconsistency
- Capture essential constraints, but don't introduce unnecessary restrictions
- Use your common sense
  - Warning: mechanical translation procedures given in this lecture are no substitute for your own judgment

http://ungenius.files.wordpress.com/2010/08/thehomer.jpg